



Abdul Wali Khan University Mardan

Department of Computer Science



BS Final Year Project Thesis

BSCS

(Hons)/Artificial

Intelligence/Software/Cyber

Security

AI Virtual Doctor: A Multi-Modal Integrated Diagnostic Engine for 34 Clinical Conditions Using Un-Nested Fine-Tuned MobileNetV2 and Late-Fusion Explainable AI (XAI)

A Thesis Submitted in Partial Fulfillment of the Requirements for the Degree of BS Computer Science.

Department of Computer Science

Abdul Wali Khan University Mardan

By

Student Name: Shahab Uddin, Amir Salman, Yasir Ali

Registration No: AWKUM-22135693, AWKUM-22135693, AWKUM-22154931

Supervisor:

Ms. Sawera Qureshi

Year: 2026

Table of Contents

Table of Figures	6
Table of Tables	8
Declaration	9
Certificate of Approval	10
Abstract	11
Chapter 1: Introduction	12
1.1 Contextual Domain Background	12
1.2 The Public Healthcare Triage Gap in Pakistan.....	12
1.3 Computational Limits of Siloed Medical AI	13
1.4 Formal Problem Statement	14
1.5 Project Technical Objectives	14
1.6 Research Questions and Hypotheses.....	15
1.7 Scope and Functional Boundaries	16
1.8 Significance of the Research Study	16
1.9 Thesis Structural Organization	17
Chapter 2: Literature Review	18
2.1 Historical Milestones in Medical Image Computing	18
2.2 Evolution of Symbolic and Tabular Symptom Processing.....	19
2.3 Early vs. Late Data Fusion Taxonomies	20
Engineering Matrix: Multi-Modal Data Fusion Trade-Offs	21
2.4 Structural Limitations of Current Medical Solutions	22
2.5 Explainable AI (XAI) and the Trust Gap in Digital Health.....	23
Chapter 3: Methodology and Architectural Pipeline	25
3.1 Proposed Multi-Modal System Framework	25
3.2 Visual Stream Topology (Un-Nested MobileNetV2 Base)	26
3.3 Tabular Stream Topology (Regularized MLP Core).....	28
3.4 Multi-Modal Late Fusion Core Mechanics	30
3.5 Mathematics of Depthwise Separable Convolutions	32
3.6 Loss Optimization and Callback Mechanics	34
3.7 Unified Feature Alignment Map	36

Engineering Matrix: Multi-Modal Stream Tensor Lifecycle.....	36
Chapter 4: System Design and Experimental Setup	37
4.1 Large-Scale Heterogeneous Dataset Profile.....	37
4.2 Proposed Multi-Model Pipeline Framework.....	38
4.3 Data Normalization and Ingestion Pipeline	39
4.4 Hardware and Software Environment Configurations.....	40
Engineering Matrix: Core Software Dependency Framework.....	41
4.5 Evaluation Metrics and Holdout Validation Strategy	42
Chapter 5: Technical Implementation & Codebase Architecture	44
5.1 Deep-Nested Multi-Wing Ingestion Engine.....	44
5.2 Unified Multi-Modal TensorFlow Dataset Pipeline.....	46
5.3 Functional Graph Compilation and Training Mechanics.....	48
5.4 Production Callbacks Implementation Script	50
5.5 Explainable AI Grad-CAM Visual Heatmap Rendering Engine.....	51
Chapter 6: Results and Discussion.....	54
6.1 Multi-Modal Diagnostic Performance Analysis.....	54
Engineering Performance Ledger: Macro Validation Metrics Across Clinical Wings.....	55
6.2 Ablation Analysis: Single-Modality Vision vs. Multi-Modal Late Fusion.....	56
6.3 Granular Error Analysis & Look-Alike Category Overlaps.....	57
6.4 Explainable AI Verification Metrics (Grad-CAM Audit).....	59
6.5 Training Dynamics, Losses, and Convergence Trajectories	62
Chapter 7: Conclusion.....	65
7.1 Thesis Summary and Project Achievement.....	65
7.2 Core Engineering Contributions	65
7.3 Experimental Validation Milestones.....	66
7.4 Final System Status and Deployment Readiness.....	67
Chapter 8: Future Work	69
8.1 Methodological Enhancements and Transformers Migration.....	69
8.2 Expanding Data Horizons: Audio and Genomic Ingestion Channels	71
8.3 Advancing Explainable AI: SHAP and Counterfactual Interventions.....	72
8.4 Decentralized Federated Learning and Edge Deployments	74
References.....	76
Appendices:.....	79

Appendix A: Comprehensive 34-Class Target Mapping Registry..... 79

 A.1 Radiology Wing Class Formats (60,000 Images) 79

 A.2 Dermatology Wing Class Formats (24,000 Images) 80

 A.3 Sonography Wing Class Formats (18,000 Images)..... 80

 A.4 Microscopic Pathology Wing Class Formats (57,000 Images) 81

Appendix B: Step-by-Step System Installation & Environment Manual 82

 B.1 Hardware Provisioning Baseline 82

 B.2 Software Installation Terminal Script..... 82

 B.3 Data Stream Pipeline Troubleshooting Directory..... 84

Appendix C: Comprehensive User Acceptance Testing (UAT) Ledger..... 84

Table of Figures

Figure 1: Chronological structural evolution of medical image recognition systems highlighting transition phases from rigid handcrafted transforms to multi-modal edge processing frameworks.	19
Figure 2: SmartArt architectural pipeline mapping the structural evolution of clinical symptom processing from legacy hardcoded symbolic engines to regularized neural embedding spaces.....	20
Figure 3: Structural comparison between traditional single-modality siloed image interpretation streams and the integrated multi-modal late-fusion processing architecture..	22
Figure 4: Operational dichotomy mapping the transparency gap between an un-auditable clinical black-box loop and the implemented explainable Grad-CAM pipeline.....	23
Figure 5: Architectural layout mapping the concurrent processing of high-dimensional visual matrices and sparse symbolic symptom vectors merging into a unified late-fusion classification core.	25
Figure 6: Dimensional transformation pipeline of the visual stream component, showcasing sequential feature extraction transitions from a raw input graphic to a compressed 256-dimensional spatial embedding.....	27
Figure 7: Dimensional transformation and regularization pipeline of the tabular symptom stream component, showcasing sequential processing transitions from a sparse 34-bit matrix to a dense 64-dimensional embedding.....	29
Figure 8: Architectural layout of the cross-modal fusion engine, mapping the parallel convergence of visual and tabular embeddings into a highly regularized 320-dimensional concatenated tensor pipeline.....	31
Figure 9: Mathematical operational workflow contrasting standard convolution against the decoupled filtering mechanics of depthwise separable convolution layers.....	32
Figure 10: Operational layout of the training execution pipeline, mapping the cyclical optimization loop from initial loss computation to conditional model callback triggers.	35
Figure 11: Dataset Class Distribution Chart.....	38
Figure 12: Sequential flow diagram of the automated data preprocessing pipeline, outlining tensor transformations from raw storage file inputs to normalized input matrices.	40
Figure 13: Hierarchical dataset segmentation schema detailing the 80/20 stratified holdout distribution pipeline from the master data registry.....	43
Figure 14: Complete 34-class evaluation matrix for the multi-modal diagnostic engine. True clinical target categories are mapped against system-predicted classifications across all integrated modalities (Chest X-ray, Ultrasound, Fungus, and Skin) to illustrate cross-class classification boundaries.....	55

Figure 15 Ablation Delta Matrix	57
Figure 16: Process execution layout of the Grad-CAM localization framework, showing the conversion of core tensor gradients into specific region-of-interest clinical attention masks.	60
Figure 17: Grad-CAM visual explanation overlay applied to a primary chest radiograph sample. The high-intensity thermal focal zones highlight localized areas of spatial pixel variation within the pulmonary region, verifying that the model's feature extraction layer is correctly focusing on structural lung patterns to predict thoracic conditions.	61
Figure 18: Explainable AI (XAI) attention localization map on an abdominal sonography image scan. The superimposed gradient heatmap isolates structural boundaries and echogenicity changes within the internal organ tissue, illustrating the model's capacity to interpret multi-modal pixel values across cross-sectional diagnostic views.	61
Figure 19: Post-hoc backpropagation saliency mapping generated for a high-magnification microscopic pathology slide. The jet-color attention overlay directly tracks individual cellular clusters and leukocyte distributions, validating that the deep convolutional base isolates specific microbiological anomalies rather than processing random background matrix noise.	62
Figure 20 Convergence Trajectory Graph	63
Figure 21: Multi-modal training and validation convergence characteristics over optimization epochs. Panel (a) illustrates the system categorical accuracy progression, demonstrating stable generalization approaching 94.0%. Panel (b) illustrates the corresponding	63
Figure 22: Future architectural layout mapping a cross-attention transformer core upgrade to replace late-fusion concatenation with deep contextual multi-modal token interactions.	70
Figure 23: Conceptual topology mapping the future four-channel data ingestion pipeline layout designed for multi-modal clinical decision-support systems.....	71
Figure 24: Architectural schema of the proposed future model-agnostic auditing framework, mapping inference outputs to parallel SHAP feature attribution and contrastive counterfactual explanations.....	73
Figure 25: Architectural topology of the proposed future decentralized federated learning paradigm, tracking bi-directional local gradient uploads and global parametric weight distribution.	74

Table of Tables

Table 1: Architectural layout of the Multi-Modal Late Fusion Core framework, illustrating the dimension concatenation of the 256-dimensional spatial visual vector and the 64-dimensional tabular symbolic vector into a consolidated 320-dimensional classification 21

Table 2: Multi-modal stream tensor lifecycle mapping structural execution stages, intermediate tensor dimension shapes, and core layer transformation logic..... 36

Table 3: Topology schema of the Master Registry Matrix outlining the multi-modal row alignment for input file paths, binary feature arrays, and categorical target targets..... 38

Table 4: Core software dependency framework detailing development platform environment configurations, production version controls, and primary operational roles within the system runtime execution layer. 41

Table 5: Engineering performance ledger showing macro validation evaluation metrics (Precision, Recall, and F1-Score) across targeted clinical diagnostic wings.....55

Table 6: Look-alike misclassification filter metrics showcasing true versus predicted cross-modal pathology overlaps alongside their primary visual error triggers..... 58

Table 7: Comprehensive summary of system experimental validation metrics contrasting project development targets against achieved performance indices..... 66

Table 8: Unified compliance and technical deployment checklist verifying core multi-modal architecture integration milestones..... 67

Table 9: Production hardware provisioning specifications outlining minimum execution environments and optimal deep learning deployment environments. 82

Table 10: System User Acceptance Testing (UAT) ledger cataloging test actions, execution expectations, and verification status across the integrated multi-modal software pipeline.. 85

Declaration

I hereby declare that this thesis is my own work and has not been submitted previously for any degree. All sources of information used have been properly acknowledged.

Student Name: Shahab Uddin, Amir Salman, Yasir Ali

Signature:

Date:

Certificate of Approval

This is to certify that the Final Year Project titled “AI Virtual Doctor: A Multi-Modal Integrated Diagnostic Engine for 34 Clinical Conditions Using Un-Nested Fine-Tuned MobileNetV2 and Late-Fusion Explainable AI (XAI)” has been carried out by Shahab Uddin, Amir Salman and Yasir Ali under my supervision and is approved for submission in partial fulfillment of the requirements for the BS Computer Science degree.

Supervisor: Ms Sawera Qureshi

Signature:

Date:

Abstract

The integration of artificial intelligence into clinical diagnostics has emerged as a critical pathway to alleviate the operational strain on healthcare infrastructures, particularly in resource-constrained regions experiencing a severe triage gap. This final year project presents the design, development, and extensive evaluation of "Ai Virtual Doctor," a unified, multi-modal diagnostic engine engineered to process heterogeneous medical data for the autonomous classification of 34 distinct clinical conditions. Moving beyond traditional single-modality medical models that evaluate radiographic or laboratory medium streams in total isolation, this architecture bridges three vital wings of diagnostic medicine: Radiology (X-Ray), Dermatology, and Microscopic Pathology. The core engineering framework utilizes a dual-stream deep neural network implemented in TensorFlow and Keras. The vision pipeline leverages an un-nested, fine-tuned MobileNetV2 backbone to extract high-dimensional visual feature vectors from a massive, multi-wing registry comprising 159,000 clinical graphics, while a parallel symbolic stream processes a 34-symptom tabular clinical matrix through a series of dense, regularized, and batch-normalized layers.

A specialized late-fusion concatenation core dynamically merges the high-dimensional visual representations with the tabular symptomatic embeddings, projecting them into a unified multi-modal feature space optimized via categorical cross-entropy loss. To establish clinical transparency and overcome the traditional black-box problem inherent to deep learning frameworks, the system integrates Gradient-weighted Class Activation Mapping (Grad-CAM) to provide explainable AI (XAI) visual saliency heatmaps that assist general practitioners by highlighting diagnostic focal zones. Evaluated via an 80/20 stratified validation strategy on localized data, the system successfully resolved complex index configurations across deep-nested bacterial, viral, fungal, and hematological subdirectory topologies. The experimental results demonstrate high precision, robust recall, and excellent clinical safety parameters, validating the framework's capability to operate as an accessible decision-support utility capable of optimizing triage workflows and mitigating diagnostic delays within the modern clinical ecosystem.

Chapter 1: Introduction

1.1 Contextual Domain Background

The deployment of autonomous computing models within the healthcare sector marks a major milestone in modern biomedical engineering. For decades, clinical diagnostics relied on manual interpretations or basic rule-based computer programs. While those early systems provided elementary data processing, they lacked the complex pattern recognition needed to handle real-world clinical data. The recent evolution of deep neural networks has fundamentally altered this landscape, turning machine learning algorithms into powerful tools for predictive healthcare. Modern medical computing uses deep convolutional networks to analyze complex medical imaging and structured multi-layer perceptrons to evaluate patient health records. This shifts the role of artificial intelligence from simple data storage to an active, transparent decision-support system capable of processing data with expert-level accuracy.

However, a major limitation of early clinical machine learning models was their complete reliance on single-modality pipelines. These networks were built to process only one type of data—such as a localized radiological scan, a specific skin lesion image, or a text-based symptom checklist—in total isolation. While these narrow models achieve high mathematical accuracy within rigid laboratory conditions, they are poorly suited for actual clinical workflows. In a real hospital environment, a human physician rarely makes a diagnostic decision based purely on a single lab slide or an isolated chest X-ray without cross-referencing the patient’s active clinical presentation, physical complaints, and acute symptom expressions. This fragmentation makes traditional single-vector models brittle and increases the risk of diagnostic errors. Consequently, contemporary biomedical research is shifting rapidly toward multi-modal data fusion setups that can ingest and process completely different data formats simultaneously.

1.2 The Public Healthcare Triage Gap in Pakistan

Public health infrastructures within developing countries, particularly across Pakistan, face severe structural and operational crises. This systemic strain is driven by expanding patient volumes, shrinking public health budgets, and a critical shortage of specialized medical personnel. Secondary and tertiary public hospitals are forced into a state of continuous resource overload. This leaves general practitioners and frontline health workers completely overwhelmed by massive backlogs of patients requiring rapid clinical routing and initial assessment. Because diagnostic testing networks operate in separate physical silos, a single patient must clear multiple disconnected testing queues—moving from radiological imaging centers to dermatological clinics or pathology laboratories—before their clinical profile can be synthesized. This highly fragmented testing workflow creates

long diagnostic delays and increases the risk of clinical oversight, directly leading to worse patient outcomes in fast-moving clinical scenarios.

This operational crisis is further exacerbated in rural and peri-urban regions, where access to expert radiologists, dermatologists, and pathologists is almost entirely non-existent. Patients must often travel long distances to major urban centers just to receive basic diagnostic interpretations, leading to catastrophic delays in treating manageable conditions. Frontline general practitioners working in underfunded rural clinics are forced to make high-stakes triage decisions without immediate access to specialist confirmations or advanced diagnostic software. This introduces intense cognitive fatigue and elevates the risk of diagnostic misclassification. To bridge this critical public health triage gap, there is an urgent need for lightweight, affordable, and highly integrated digital diagnostics. Deploying cross-wing, multi-modal screening tools at the primary point of care can help general practitioners rapidly filter out normal cases and identify high-risk anomalies, optimizing patient routing and maximizing resource utilization across public hospitals.

1.3 Computational Limits of Siloed Medical AI

From a software engineering and machine learning perspective, existing healthcare applications fail to resolve this triage crisis because they remain structurally disconnected. The vast majority of current computer vision models are heavily siloed within single medical specialties. For example, a deep neural network trained to detect pulmonary nodules on a chest radiograph completely lacks the capability to interpret a skin infection or analyze microscopic blood cell distributions. Forcing an underfunded clinic to deploy, maintain, and run multiple separate single-modality software instances introduces massive logistical friction, demands excessive computing power, and complicates system integration. This makes advanced machine learning tools impractical for deployment on standard clinic computers.

Furthermore, image-only classification networks display high false-positive rates because they completely lack access to the patient's acute physical symptoms. A model evaluating a dense chest area on an X-ray film cannot cross-reference the presence of an active high fever, shivering, or a productive cough. This missing context frequently causes single-modality models to misinterpret benign anatomical variations or temporary image artifacts as malignant pathologies, leading to unnecessary referrals and increased diagnostic anxiety[19]. Without a unified deep learning model that handles cross-modal feature fusion, automated clinical aids remain isolated tools. They remain detached from the comprehensive, multi-faceted reality of human clinical diagnostic procedures and are unable to survive the noisy, unpredictable data environments of real-world public healthcare systems.

1.4 Formal Problem Statement

The operational efficacy of automated primary care triage within Pakistan's public health network is severely limited by two interconnected challenges: the institutional fragmentation of multi-wing diagnostic testing workflows and the computational limitations of single-modality "black-box" machine learning pipelines. Underfunded primary clinics lack a unified, resource-optimized software utility capable of simultaneously processing disparate medical data streams, such as high-contrast chest X-rays, surface tissue dermatological graphics, and microscopic pathology slides. This lack of integration forces clinicians to rely on slow, disconnected testing loops, causing severe diagnostic delays and increasing triage errors in high-volume public hospitals.

At the same time, existing clinical machine learning applications fail to provide necessary operational safety because they lack contextual awareness and structural transparency. Image-only classification models generate high-variance errors and elevated false-positive rates because they cannot cross-reference visual patterns with a patient's self-reported physical symptoms. Furthermore, the opaque nature of traditional deep neural networks prevents general practitioners from verifying the underlying reasoning behind a model's multi-class prediction. Without a unified, lightweight, dual-stream architecture that performs late-fusion vector concatenation to merge visual tensors with tabular symptom matrices—and overlays explainable visual heatmaps to show the model's focus—automated medical aids remain untrusted, impractical, and isolated from the multi-dimensional requirements of frontline clinical triage operations.

1.5 Project Technical Objectives

The overarching engineering objective of this final year project is to develop, optimize, and validate a unified multi-modal deep learning software system that performs concurrent, cross-wing primary care diagnostic triage. To achieve this broad goal with strict architectural and scientific rigor, the development lifecycle is mapped across five discrete, quantifiable technical objectives. The first milestone requires building an automated, path-aware data ingestion and preprocessing pipeline capable of recursively parsing 159,000 highly diverse clinical image assets. This pipeline must resolve nested local directory topologies, bypass inconsistent zip folder layers, correct naming shorthand prefixes, and standardize variable graphic structures into a uniform $224 \times 224 \times 3$ floating-point tensor array ready for model input.

The second core objective focuses on engineering a parallel, dual-stream neural network architecture within the TensorFlow and Keras environments. This requires setting up a vision-processing stream using an un-nested, fine-tuned MobileNetV2 base with frozen pretrained ImageNet parameters for spatial extraction, alongside a symbolic multi-layer

perceptron stream optimized to transform a sparse 34-symptom binary patient matrix into continuous continuous embeddings. The third objective involves designing a dedicated late-fusion concatenation core that joins the independent visual and symbolic representation streams along a single axis, mapping complex cross-modal patterns without experiencing geometric imbalance or gradient disruption. The fourth milestone is to achieve multi-class diagnostic differentiation across 34 mutually exclusive target conditions spanning Radiology, Dermatology, and Microscopic Pathology. The final objective requires integrating an explainable AI rendering core using Gradient-weighted Class Activation Mapping (Grad-CAM) to generate visual saliency heatmaps, transforming a traditional "black-box" model into a transparent, auditable decision-support utility for frontline general practitioners.

1.6 Research Questions and Hypotheses

The structural design and experimental setup of the "Ai Virtual Doctor" multi-modal architecture are organized to investigate specific engineering trade-offs and clinical deployment paradigms. This study seeks to systematically resolve three primary research questions. First, how effectively can a lightweight, mobile-optimized deep convolutional neural network backbone, specifically an un-nested MobileNetV2, extract generalized spatial features across completely distinct clinical imaging modalities—such as high-contrast chest radiographs, surface dermatological lesions, and microscopic fluid smear slides—while operating within a restricted local parameter footprint? Second, to what degree does the mathematical concatenation of a low-dimensional tabular symptom embedding layer with a high-dimensional visual feature vector at the late-fusion layer minimize cross-class misclassification errors and false-positive rates compared to traditional single-modality image-only pipelines? Third, can the integration of gradient-based visual saliency maps successfully establish clinical trust and transparency for general practitioners running multi-class triage inferences in underfunded environments?

To address these questions, the experimental pipeline evaluates two core research hypotheses. The primary hypothesis (H_1) states that the late-fusion integration of structured patient symptom bit-vectors with dense image feature maps will yield a statistically significant improvement in multi-class classification metrics (Precision, Recall, and F1-Score) when distinguishing between visually look-alike conditions, compared to an isolated single-modality vision network. The secondary hypothesis (H_2) asserts that an un-nested, fine-tuned MobileNetV2 backbone will achieve top-1 validation metrics comparable to heavy, deep convolutional ensemble networks while reducing overall floating-point operations per second (FLOPs) and computational parameter weights. This reduction allows the multi-wing model to execute rapid real-time diagnostic screening on consumer-grade computers lacking specialized GPU infrastructure.

1.7 Scope and Functional Boundaries

The operational, architectural, and experimental boundaries of this project are strictly defined to ensure deep engineering quality, mathematical stability, and realistic deployment scopes within the timeline of a senior undergraduate thesis. Architecturally, the multi-wing vision pipeline is limited to exactly three clinical testing specialties to ensure varied data inputs while preventing uncontrolled training drift. The radiological wing is exclusively confined to chest X-ray views focusing on fifteen target thoracic conditions, including asthma, atelectasis, COVID-19, cardiomegaly, consolidation, edema, effusion, emphysema, fibrosis, infiltration, mass, nodule, pneumonia, pneumothorax, and tuberculosis. The dermatological wing is strictly bounded to acute surface tissue conditions, categorizing localized bacterial, fungal, parasitic, and viral skin diseases. The microscopic pathology wing is limited to hematological and microbiological slides, classifying specific leukocyte distributions, thrombocytes, malaria variants, and common bacterial pathogens such as staphylococcus and streptococcus.

On the symbolic data processing side, the clinical input layer is explicitly restricted to a fixed 34-symptom binary matrix, which captures key physiological flags such as high fever, chills, jaundice, chest pain, and a productive cough. This choice keeps execution fast and prevents memory bottlenecks. The project explicitly excludes unstructured textual logs, conversational voice data, or streaming genomic arrays. From a software perspective, the implementation is built using Python 3.12, TensorFlow 2.16, and Keras 3, utilizing categorical cross-entropy loss over a fixed target label set of 34 classes. Finally, the deployment scope is framed as an explainable, interactive decision-support utility meant for a general practitioner's terminal. It is explicitly not designed to operate as an unguided, fully autonomous diagnostic tool, keeping a clear line between automated medical triage and final physician authorization.

1.8 Significance of the Research Study

The technical and social significance of this research lies in its ability to provide a scalable, low-overhead software solution tailored to address the severe healthcare resource crises in developing public health networks, such as those across Pakistan. By combining radiology, dermatology, and microscopic pathology into a single deep learning instance, the "Ai Virtual Doctor" system significantly reduces the compute footprint, memory demands, and system integration friction that typically prevent advanced clinical AI from being deployed on standard clinic computers. This multi-modal fusion architecture helps eliminate the high-variance failures common in single-modality image networks. This feature provides frontline healthcare workers with exceptional diagnostic reliability, even when processing noisy, low-contrast, or imperfect real-world scans.

Furthermore, integrating explainable AI through a customized Grad-CAM pipeline directly resolves the traditional "black box" trust problem that often hinders the real-world adoption of deep learning models in medicine. Instead of presenting an isolated prediction percentage, the framework generates real-time visual attention maps directly over the medical graphics. This allows general practitioners to instantly cross-verify the AI's diagnostic focus against established clinical landmarks. In rural primary care clinics, this capability allows frontline workers to rapidly filter out normal cases and identify high-risk anomalies. This shifts the healthcare delivery model from a slow, fragmented testing workflow to an efficient, integrated triage system that optimizes patient routing and shortens diagnostic loops.

1.9 Thesis Structural Organization

This thesis document is structured into seven distinct, highly comprehensive chapters that systematically track the development of this multi-modal final year project from initial concept down to final validation. Chapter 1 introduces the study, outlining the conceptual background, the public healthcare triage gap in Pakistan, the computational limits of siloed models, the formal problem statement, project objectives, hypotheses, and scope parameters. Chapter 2 provides an exhaustive Literature Review, analyzing historical milestones in medical image computing, the evolution of clinical symptom processing models, the mathematics behind early and late data fusion architectures, and the current research gaps in explainable medical machine learning. Chapter 3 explains the Methodology and Architectural Pipeline, detailing the dual-stream functional network layout, depthwise separable convolutions, un-nested MobileNetV2 tuning, and late-fusion concatenation layers.

Chapter 4 details the System Design and Experimental Setup, presenting the 159,000-image dataset profile, master registry indexing configurations, data normalization methods, hardware/software environment parameters, and stratified validation strategies. Chapter 5 focuses on the Technical Implementation, showcasing the compiled training loops, path-aware directory parsing modules, custom `tf.data` dataset generator scripts, and Grad-CAM activation mapping code blocks. Chapter 6 provides an in-depth Results and Discussion analysis, evaluating the multi-class performance report across 34 target classes, breaking down cross-classification errors via a comprehensive confusion matrix, and conducting ablation reviews. Finally, Chapter 7 concludes the thesis by summarizing key engineering accomplishments, discussing current system limitations, and suggesting future research enhancements.

Chapter 2: Literature Review

2.1 Historical Milestones in Medical Image Computing

The evolution of medical image computing has transitioned through several distinct technological generations, moving from basic edge-detection algorithms to deep, self-optimizing feature extractors. In the early eras of digital health systems, automated computer-aided diagnosis (CAD) relied entirely on handcrafted spatial filters and rigid mathematical transforms. Engineers used techniques like the Hough Transform to locate circular contours in bone radiographs, or applied Gabor filters to map surface tissue variations in dermatological imagery. These early systems were heavily constrained because they were entirely non-adaptive. If a clinical scan exhibited slight variations in exposure, structural orientation, or localized artifact noise, the handcrafted mathematical filters failed immediately, resulting in high false-negative rates in complex hospital environments.

The introduction of Deep Convolutional Neural Networks (CNNs) completely altered this paradigm by shifting the burden of feature extraction from human engineers to automated backpropagation weight learning. Early architectures proved that stacked convolutional layers could automatically extract hierarchically organized visual features directly from raw pixels—moving from low-level edges and corners in shallow layers to complex pathological structures in deeper dimensions.

However, applying these heavy computer vision networks directly to healthcare presented severe operational challenges. Models like AlexNet or VGG16 contained tens of millions of trainable parameters, demanding massive memory bandwidth and specialized GPU clusters. This computational overhead made early deep learning applications completely impractical for resource-constrained primary care clinics, highlighting the critical research need for highly optimized, mobile-adapted convolutional structures.




Figure 1: Chronological structural evolution of medical image recognition systems highlighting transition phases from rigid handcrafted transforms to multi-modal edge processing frameworks.

2.2 Evolution of Symbolic and Tabular Symptom Processing

Parallel to the advancements in computer vision, the computational processing of patient-reported symptoms and tabular clinical histories progressed along a separate developmental path. Early medical artificial intelligence applications used symbolic expert systems, such as MYCIN, which relied on massive databases of hardcoded, human-engineered "if-then-else" logic rules to infer infectious disease profiles. While these early rule-based setups provided clear diagnostic transparency, they were completely incapable of handling missing values, noisy inputs, or the non-linear symptom correlations that occur

in real-world clinical practices. If a patient presented with an atypical variation of a common condition, the rigid rule-based engines failed to compute an accurate match, rendering them brittle and impractical for dynamic clinical environments.




Figure 2: SmartArt architectural pipeline mapping the structural evolution of clinical symptom processing from legacy hardcoded symbolic engines to regularized neural embedding spaces.

To resolve these structural limitations, modern biostatistical computing transitioned toward machine learning algorithms capable of processing tabular matrices as continuous vector spaces. Researchers deployed architectures like Multi-Layer Perceptrons (MLPs), Support Vector Machines (SVMs), and Gradient Boosted Decision Trees to analyze patient diagnostic indicators.

When applied to sparse symptom matrices, an optimized MLP maps true-or-false clinical signs into low-dimensional continuous vector embeddings. By adding regularization layers—such as batch normalization vectors and dropout masks—modern tabular networks extract deep mathematical relationships between distinct physical complaints, such as correlating high fever and shivering with specific infection types. This architectural shift allows symbolic models to maintain high predictive stability, even when processing highly sparse or incomplete patient data profiles.

2.3 Early vs. Late Data Fusion Taxonomies

In multi-modal machine learning engineering, data fusion methods are divided into three primary categories based on where the independent data streams are integrated within the

execution graph: Early Fusion (input-level), Intermediate Fusion (layer-level), and Late Fusion (decision-level). Early fusion architectures combine different input sources, such as raw image pixels and tabular text vectors, directly at the initial input layer before any major feature extraction occurs. While this strategy lets the network analyze raw multi-modal correlations from the very start of training, it introduces severe geometric alignment errors when processing highly disparate data formats. Attempting to concatenate high-dimensional, dense image matrices directly with sparse, low-dimensional tabular bit-vectors creates massive mathematical imbalances that destabilize gradient calculations and cause the network to ignore the smaller tabular data stream entirely.

Engineering Matrix: Multi-Modal Data Fusion Trade-Offs

Table 1: Architectural layout of the Multi-Modal Late Fusion Core framework, illustrating the dimension concatenation of the 256-dimensional spatial visual vector and the 64-dimensional tabular symbolic vector into a consolidated 320-dimensional classification

Fusion Strategy	Integration Point	Architectural Complexity	Primary Failure Mode	Best Use Case
Early Fusion	Input Layer	Low	Input Matrix Dimension Imbalances; Gradient Destabilization	Homogeneous Data Streams (e.g., Dual Video Feeds)
Intermediate Fusion	Hidden Layers	High	Extreme Architectural Overhead; High Training Latency	Deeply Correlated Multi-Sensor Robotics Data
Late Fusion	Concatenation Core	Balanced	Fails if Latent Feature Vectors Are Not Density-Aligned	Heterogeneous Systems (e.g., Medical Scans + Bit-Vectors)

Late fusion architectures solve this geometric conflict by keeping the independent data streams entirely separate throughout the primary feature extraction process. Each data format is handled by its own specialized neural network stream, allowing the visual pipeline to focus exclusively on spatial details and the tabular pipeline to focus on symbolic pattern matching. The independent latent feature vectors are only merged at the final layers

of the execution graph via a structured concatenation core. Fusing the data at this stage ensures that both modalities are compressed into dense, optimized representations before they interact, preventing either stream from dominating the network. This structural separation maintains clean gradient transitions, protects the model from input-level scale imbalances, and ensures balanced multi-class training performance across the entire dataset.

2.4 Structural Limitations of Current Medical Solutions

Despite the rapid development of deep learning models for healthcare, existing software applications remain heavily limited by two major structural flaws: a complete lack of multi-wing adaptability and high computational resource demands. Most contemporary medical image models are strictly single-purpose tools built for isolated diagnostic tasks. A network trained to classify chest X-rays cannot analyze dermatological lesions or evaluate microscopic pathology slides. Forcing an underfunded community clinic to deploy, host, and maintain multiple independent single-modality AI software tools introduces massive system integration friction, demands excessive computing power, and complicates daily workflows, making advanced clinical software impractical for standard clinic computers.




Figure 3: Structural comparison between traditional single-modality siloed image interpretation streams and the integrated multi-modal late-fusion processing architecture.

Furthermore, current image-only models suffer from elevated false-positive rates because they completely lack access to the patient's active physical symptoms. For example, an image-only vision network evaluating a dense pulmonary zone on a chest radiograph cannot cross-reference the presence of an active high fever or a productive cough. This missing context frequently causes single-modality models to misinterpret benign anatomical variations or temporary image noise as malignant conditions, leading to unnecessary specialist referrals and increased diagnostic anxiety. Without a unified deep learning framework that combines visual spatial processing with symbolic symptom embeddings, automated medical aids remain isolated tools that are poorly suited to handle the unpredictable data environments of real-world public healthcare triage systems.

2.5 Explainable AI (XAI) and the Trust Gap in Digital Health




Figure 4: Operational dichotomy mapping the transparency gap between an un-auditable clinical black-box loop and the implemented explainable Grad-CAM pipeline.

To bridge this critical trust gap, contemporary biomedical engineering places a high priority on integrating Explainable AI (XAI) utilities directly into model pipelines. Techniques like Gradient-weighted Class Activation Mapping (Grad-CAM) solve this opacity issue by using the gradient vectors flowing into the final convolutional layer to calculate spatial importance maps^[4]. This process generates colorized visual heatmaps that highlight the exact regions within a medical image that drove the model's final classification.

By explicitly showing where the network is focusing—such as isolating active tissue lesion boundaries or highlighting localized pulmonary anomalies—the system transforms from an un-auditable prediction engine into a transparent, collaborative decision-support tool. This capability allows general practitioners to quickly verify the AI's reasoning against established clinical landmarks, ensuring high operational safety during frontline triage workflows.

Chapter 3: Methodology and Architectural Pipeline

3.1 Proposed Multi-Modal System Framework

The architectural core of the "Ai Virtual Doctor" engine is designed as an integrated, multi-modal deep learning platform that fuses high-dimensional visual tensors with low-dimensional symbolic matrices. Traditional single-modality networks deployed within public healthcare settings create operational friction by requiring separate instances for processing distinct medical data forms. To circumvent these runtime inefficiencies, the proposed framework unifies three major diagnostic wings—Radiology, Dermatology, and Microscopic Pathology—into a single, consolidated model graph. The system achieves this by running a dual-input pipeline that decouples initial feature extraction stages based on input formats, processing clinical graphics and discrete patient-reported symptoms independently before executing late-fusion merging.




Figure 5: Architectural layout mapping the concurrent processing of high-dimensional visual matrices and sparse symbolic symptom vectors merging into a unified late-fusion classification core.

The underlying software architecture is implemented within the TensorFlow 2.16 and Keras 3 ecosystems, establishing an asynchronous, end-to-end dataflow from raw file ingestion down to multi-class prediction probabilities. When a patient query is initialized, the system passes the localized file path of the diagnostic scan and a 34-dimensional binary symptom vector to their respective parallel stream entry blocks. The vision-processing stream extracts complex structural and pathological features from the images, while the symbolic processing stream compresses sparse clinical signs into dense embedding spaces. These independent vector streams meet at a specialized vector concatenation core that links them into a single cross-modal tensor space. This combined feature map is then passed through regularized dense classification layers to map non-linear symptom-image correlations, generating a final stable softmax probability distribution across 34 mutually exclusive target conditions.

3.2 Visual Stream Topology (Un-Nested MobileNetV2 Base)

The visual processing stream of the model is designed to perform high-efficiency spatial feature extraction while operating within the restricted parameter and memory boundaries of local clinic computers. To achieve a balance between high multi-class accuracy and small computational footprint, the architecture uses an un-nested MobileNetV2 backbone pretrained on ImageNet feature registries. The top classification layer of the native MobileNetV2 network is discarded, opening the core feature extraction base to accept input tensors resized to a standardized format of $224 \times 224 \times \text{pixels}$. This base relies on depthwise separable convolutions that break standard convolutional operations down into a single spatial filter followed by a pointwise 1×1 channel mapping. This technique reduces overall parameter dependencies and floating-point operations per second (FLOPs) by nearly nine-fold compared to traditional convolutional architectures.




Figure 6: Dimensional transformation pipeline of the visual stream component, showcasing sequential feature extraction transitions from a raw input graphic to a compressed 256-dimensional spatial embedding.

To protect the model from distorting the specialized structural patterns learned during pretraining, all foundational convolutional layers inside the MobileNetV2 base are locked (`base_model.trainable = False`), ensuring their weights remain non-trainable during early optimization runs. Input graphics have their pixel intensity arrays normalized to a continuous scale between 0.0 and 1.0.

As an input tensor travels through the inverted residual blocks and linear bottlenecks of the frozen base, it is transformed from a flat pixel grid into a dense $7 \times 7 \times 1280$ feature map. This high-dimensional array is passed directly to a Global Average Pooling 2D (GAP) layer. The GAP layer computes the mathematical mean of each channel matrix, converting the 3D feature tensor into a compact, 1280-dimensional visual embedding vector that represents the core spatial characteristics of the clinical scan. This vector is then compressed through a dense projection layer down to 256 dimensions to ensure geometric alignment prior to data fusion.

3.3 Tabular Stream Topology (Regularized MLP Core)

Operating in parallel with the computer vision stream, the symbolic processing pipeline is engineered to ingest, parse, and encode structured patient clinical indications, turning a patient's physical symptoms into a continuous vector embedding space. The input interface consists of a dedicated dense layer that receives a fixed 34-dimensional binary matrix from the user terminal, where each discrete bit signifies the absolute presence (1.0) or absence (0.0) of specific clinical indicators, including high fever, chills, jaundice, shivering, chest pain, and a productive cough. Because these binary symptom vectors are inherently sparse, passing them directly into an unregularized deep neural network can cause extreme gradient variance and severe model overfitting. To address this risk, the tabular stream implements a Multi-Layer Perceptron (MLP) architecture configured with active regularization layers.




Figure 7: Dimensional transformation and regularization pipeline of the tabular symptom stream component, showcasing sequential processing transitions from a sparse 34-bit matrix to a dense 64-dimensional embedding.

The incoming 34-bit symptom array is passed directly into an initial Dense layer containing 64 hidden units powered by the Rectified Linear Unit (ReLU [Type equation here.](#)) activation function. To stabilize internal covariate shifts and accelerate optimization runs across variable training batch sizes, a Batch Normalization layer is applied immediately after this initial dense projection. This is followed by a specialized Dropout layer with a rate of 0.2, which randomly deactivates 20% of the active hidden units during each training step. This design breaks artificial co-dependencies among neurons and forces the network to learn generalized symptom combinations. The filtered signals are then routed through a second Dense layer with 64 units and ReLU activation, producing a dense, 64-dimensional symbolic feature vector that represents the patient's symptomatic state.

3.4 Multi-Modal Late Fusion Core Mechanics

The core innovation of the "Ai Virtual Doctor" architecture lies in its late-fusion core, which integrates the independent feature representations extracted from the visual and tabular streams into a unified diagnostic layer. Traditional early-fusion techniques merge data streams at the raw input layer, which often fails in heterogeneous clinical models because the massive mathematical dimensionality of an image tensor completely overwhelms compact tabular vectors, rendering the symptom data useless. The late-fusion design avoids this issue by allowing the visual and symbolic processing streams to run independently through their respective deep pipelines, extracting high-level feature vectors before performing cross-modal combination.




Figure 8: Architectural layout of the cross-modal fusion engine, mapping the parallel convergence of visual and tabular embeddings into a highly regularized 320-dimensional concatenated tensor pipeline.

The fusion core uses a vector concatenation layer to join the compressed visual feature vector with the symbolic symptom embedding, producing a single, multi-modal feature tensor. This combined tensor is passed into a fully connected Dense layer with 128 hidden units and ReLU activation, which projects the fused data into a shared multi-modal embedding space to model complex, cross-modal relationships. To maintain training stability and prevent the larger image features from dominating the network weights, this layer is immediately followed by a Batch Normalization layer and a final Dropout layer with a rate of 0.3. This highly regularized structure ensures that the network successfully

models complex, non-linear relationships between visual clinical anomalies and physical symptoms without overfitting.

3.5 Mathematics of Depthwise Separable Convolutions

The computational efficiency that allows the visual backbone to operate smoothly on standard, non-GPU local clinic computers is driven by the mathematical structure of depthwise separable convolutions. Traditional convolutional layers process spatial dimensions and channel dependencies simultaneously within a single step. For an input tensor of dimensions $D_F \times D_F \times M$ mapped by a kernel matrix of size $D_K \times D_K \times M \times N$, a standard convolution produces a computational complexity cost defined by:

$$C_{STD} = D_F^2 \cdot M \cdot N \cdot D_K^2$$

This simultaneous processing creates an intense computational burden and a large parameter footprint, causing standard CNNs to drop significantly in inference speeds when deployed without expensive GPU hardware.




Figure 9: Mathematical operational workflow contrasting standard convolution against the decoupled filtering mechanics of depthwise separable convolution layers.

STANDARD CONVOLUTION (Simultaneous Spatial & Channel Filtering):

[Input: $D_f \times D_f \times M$] $\Rightarrow$ (Kernel: $D_k \times D_k \times M \times N$) $\Rightarrow$ [Output: $D_f \times D_f \times N$]

* Total Multiply-Accumulate Operations = $D_f^2 \cdot M \cdot N \cdot D_k^2$

DEPTHWISE SEPARABLE CONVOLUTION (Decoupled Filtering Process):

Step 1: Depthwise Spatial Filtering

[Input: $D_f \times D_f \times M$] $\rightarrow$ (Kernel: $D_k \times D_k \times 1 \times M$) $\rightarrow$ [Intermediate: $D_f \times D_f \times M$]

* Step 1 Operations = $D_f^2 \cdot M \cdot D_k^2$

Step 2: Pointwise Channel Mapping

[Intermediate: $D_f \times D_f \times M$] $\rightarrow$ (Kernel: $1 \times 1 \times M \times N$) $\rightarrow$ [Output: $D_f \times D_f \times N$]

* Step 2 Operations = $D_f^2 \cdot M \cdot N$

$$\text{Combined Cost} = (D_f^2 \cdot M \cdot D_k^2) + (D_f^2 \cdot M \cdot N)$$

Depthwise separable convolutions solve this bottleneck by breaking the convolutional operation down into two separate, specialized layers: a depthwise convolution and a pointwise convolution. The depthwise convolution applies a single spatial filter per input channel, mapping spatial patterns across the tensor without changing the overall channel depth. This operation yields a localized computational cost of $D_F^2 \cdot M \cdot D_K^2$.

Immediately following this, the pointwise convolution applies a 1×1 convolution across all channels, projecting the spatial features into a new channel depth of N . The computational cost for the pointwise stage is defined as $D_F^2 \cdot M \cdot N$. By combining these two distinct operations, the total computational complexity of the visual backbone drops to:

$$C_{\text{separable}} = (D_F^2 \cdot M \cdot D_K^2) + (D_F^2 \cdot M \cdot N)$$

When evaluating the efficiency gains by comparing the two computational profiles, the structural reduction factor is expressed as:

$$R = \frac{(D_F^2 \cdot M \cdot D_K^2) + (D_F^2 \cdot M \cdot N)}{D_F^2 \cdot M \cdot N \cdot D_K^2} = \frac{1}{N} + \frac{1}{D_K^2}$$

Since the visual backbone uses a core kernel size where $D_K = 3$, depthwise separable convolutions reduce the computational workload and parameter count by approximately eight to nine times compared to standard convolutions. This mathematical optimization preserves high top-1 classification accuracy while keeping the system lightweight enough for rapid primary care deployments.

3.6 Loss Optimization and Callback Mechanics

The optimization pipeline of the multi-modal engine is built to handle stable multi-class differentiation across 34 target categories by minimizing an objective loss function over a heavily regularized parameter space. Because the project involves a mutually exclusive multi-class classification task where each patient profile maps to a single primary clinical condition, the system uses Categorical Cross-Entropy loss. This loss function measures the information distance between the model's predicted softmax probability vector and the true, one-hot encoded label matrix. It is mathematically defined as:

$$L_{CCE} = - \sum_{i=1}^N \sum_{j=1}^C y_{i,j} \log(\hat{y}_{i,j})$$

where N represents the operational batch size, C represents the 34 target classes, $y_{i,j}$ is the ground-truth binary indicator, and $\hat{y}_{i,j}$ is the predicted softmax probability.




Figure 10: Operational layout of the training execution pipeline, mapping the cyclical optimization loop from initial loss computation to conditional model callback triggers.

The network weights are updated using the Adam optimizer, initialized with a learning rate of 1×10^{-3} , which dynamically adjusts individual parameter step sizes based on running updates of the first and second moments of the gradients. To prevent training stagnation and protect the model from getting stuck in local minima, the optimization pipeline integrates three automated Keras callbacks.

A `ModelCheckpoint` utility continuously monitors validation accuracy and saves the best model state (`multimodal_model_best.keras`). Concurrently, an `EarlyStopping` callback monitors validation loss and halts training if the metric fails to improve for 3 consecutive epochs, restoring the best weights. Finally, a `ReduceLROnPlateau` callback dynamically drops the learning rate by a factor of 0.2 if validation loss plateaus for 2 consecutive epochs, ensuring smooth and stable convergence.

3.7 Unified Feature Alignment Map

To guarantee that cross-modal concatenation layer execution does not cause shape mismatches or model compiling errors, the internal matrix shapes of both input channels are explicitly tracked, transformed, and locked at every step of the computational graph. Visual data shapes are tracked through deep spatial filters, pooling grids, and flat dense layers. Concurrently, the tabular symptoms are expanded through linear layers to ensure density balance before merging.

Engineering Matrix: Multi-Modal Stream Tensor Lifecycle

Table 2: Multi-modal stream tensor lifecycle mapping structural execution stages, intermediate tensor dimension shapes, and core layer transformation logic.

Execution Stage	Visual Stream Shape	Tabular Stream Shape	Core Transformation Logic
Data Input	(None, 224, 224, 3)	(None, 34)	Raw Ingestion & Synchronized Tensor Formatting
Feature Extraction	(None, 7, 7, 1280)	(None, 64)	Depthwise Spatial Filtering vs. MLP Dense Mapping
Pooling & Pooling Output	(None, 1280)	(None, 64)	Global Average Pooling 2D Channel Reduction
Pre-Fusion Alignment	(None, 256)	(None, 64)	Dense Compression Layer for Geometric Balance
Late Fusion Core	Merged Vector Space: (None, 320)	Cross-Modal Concat Vector Concatenation Layer	
Output Classification	Unified Target Prediction: (None, 34)	Softmax Layer Multi-Class Normalization Vector	

This alignment map ensures that the 256-dimensional visual feature vector and the 64-dimensional tabular symptom embedding are combined into a perfectly proportioned 320-dimensional tensor. By managing tensor dimensions throughout the execution graph, the framework avoids shape mismatches and gradient drops during backpropagation. This mathematical optimization guarantees stable multi-class validation performance across all 34 target clinical categories.

Chapter 4: System Design and Experimental Setup

4.1 Large-Scale Heterogeneous Dataset Profile

The empirical verification of the multi-modal classification system is grounded in a large-scale, heterogeneous clinical dataset comprising 159,000 high-resolution medical graphics paired with structured clinical reports[21]. To ensure comprehensive diagnostic coverage across diverse healthcare wings, the dataset is systematically organized into three distinct operational categories: Radiology, Dermatology, and Microscopic Pathology.

The Radiology wing accounts for 55,000 chest radiographs focusing on major thoracic conditions. The Dermatology wing contains 52,000 surface tissue graphics mapping localized skin infections. The Pathology wing consists of 52,000 high-magnification microscopic fluid and smear slides detailing cellular components and bacterial pathogen distributions.

This large-scale image collection is managed via a centralized database ledger, `Master_Clinical_Reports_Realistic.csv`. This master registry file serves as the core database matrix for the system, mapping each unique image filename to its corresponding 34-symptom binary vector, its clinical wing assignment, and its diagnostic target category.

During initialization, an automated disk-scraping module recursively maps the local file system to bypass deep-nested folder layers and resolve path inconsistencies across different storage directories. This step pairs each file path with its exact row entries in the master registry. This automated data mapping ensures that the model inputs remain completely synchronized, preventing matrix alignment errors and protecting the data pipeline from runtime execution failures.




Figure 11: Dataset Class Distribution Chart

4.2 Proposed Multi-Model Pipeline Framework

To resolve the structural limitations of single-modality medical applications, this study introduces a unified engineering architecture that coordinates the processing of independent feature spaces before classification execution. The master registry uses a structured, index-aligned format where each row represents a complete patient diagnostic

record. The columns are strictly divided into three functional segments: file management track

Table 3: Topology schema of the Master Registry Matrix outlining the multi-modal row alignment for input file paths, binary feature arrays, and categorical target targets.

Dataset Index	File Identification Reference Path	34-Symptom Tabular Binary Vector Array	Evaluated Ground-Truth Target Class
001	"/dataset/radiology/xray_091.png"	[1.0, 0.0, 0.0, 0.0, ..., 1.0]	COVID-19 Acute Infiltration
002	"/dataset/dermatology/lesion_52.jpg"	[0.0, 1.0, 0.0, 0.0, ..., 0.0]	Bacterial Skin Infection
003	"/dataset/pathology/smear_104.jpeg"	[0.0, 0.0, 1.0, 0.0, ..., 0.0]	Leukocyte Shift (Neutrophil)

ing, the 34-symptom binary checklist array, and the target diagnostic text labels.

The symptom checklist array is stored as a series of 34 binary columns, where each discrete feature maps directly to a specific physiological sign or patient complaint, such as high fever, jaundice, or chest pain. These columns are processed as a continuous bit-vector by the Multi-Layer Perceptron (MLP) stream. The target diagnostic labels are stored as string values representing the 34 mutually exclusive target conditions. During data loading, these text labels undergo a label encoding transformation that maps each string category to a sequential integer between 0 and 33. These integers are then converted into 34-dimensional binary target vectors using a one-hot encoding function, formatting the target outputs to match the model's final softmax activation layer.

4.3 Data Normalization and Ingestion Pipeline

To ensure stable gradient descent and prevent input scaling variations from causing model divergence during training, all raw data streams pass through an automated normalization pipeline. For the computer vision pipeline, raw medical graphics are ingested via TensorFlow's native I/O modules, which decode variable byte streams into 3-channel RGB pixel matrices. An automated resizing transformation standardizes all variable image dimensions to a uniform target resolution of 224 × 224 pixels. These integer pixel values (ranging from 0 to 255) are converted to 32-bit floating-point numbers and divided by 255.0, scaling all input image features to a standard range between 0.0 and 1.0.

Simultaneously, the tabular data pipeline processes the 34 binary symptom columns from the master registry, ensuring they are cast as 32-bit floating-point arrays to match the numerical precision of the image features. The text-based diagnostic labels undergo a label encoding transformation that maps each string category to a clean sequential integer between 0 and 33. These integer indices are then processed using a one-hot encoding function, converting them into 34-dimensional binary target vectors. This dual-stream preprocessing setup converts disparate file types into standardized, synchronized tensor streams, feeding clean matrices into the model's parallel input layers.




Figure 12: Sequential flow diagram of the automated data preprocessing pipeline, outlining tensor transformations from raw storage file inputs to normalized input matrices.

4.4 Hardware and Software Environment Configurations

The training, optimization, and evaluation of the framework were conducted within a dedicated deep learning environment configured to ensure stable execution and reproducible results. The hardware setup consisted of an AMD Ryzen multicore processor

paired with an NVIDIA GeForce RTX graphics processing unit to accelerate convolutional operations via CUDA and cuDNN libraries. The system was equipped with 16 GB of high-speed RAM to handle the large data arrays generated when caching data pipelines for the 159,000-image registry.

Engineering Matrix: Core Software Dependency Framework

Table 4: Core software dependency framework detailing development platform environment configurations, production version controls, and primary operational roles within the system runtime execution layer.

Software Dependency	Production Version	Primary Operational Role
Operating System	Microsoft Windows 11	Host Environment & Thread Scheduling Core
Python Core Runtime	Version 3.12.2	Ingestion Scripting & In-Memory Execution
TensorFlow Engine	Version 2.16.1	Static Computational Graph Architecture
Keras API Wrapper	Version 3.10.0	Functional Stream Object Graph Layer
Pandas Library	Version 2.2.1	CSV Ingestion & Subdirectory Dataframe Ledger
NumPy Array Framework	Version 1.26.4	Vector Transformations & Matrix Manipulations
Scikit-Learn Utility	Version 1.4.1	Holdout Stratification & Validation Scoring

On the software side, the environment was built on a 64-bit Windows 11 operating system running Python 3.12. The core neural network operations were managed using TensorFlow 2.16.1 and Keras 3.10.0, utilizing the TensorFlow backend for tensor graphs[5][6]. Data manipulation, file indexing, and matrix transformations were handled using Pandas 2.2 and NumPy 1.26[7][8]. Mathematical evaluation metrics, classification reports, and confusion matrices were computed using Scikit-Learn 1.4[9]. Visual charts and saliency maps were generated using Matplotlib 3.8 and Seaborn 0.13[11]. System logging filters were set to

suppress unnecessary background runtime warnings, keeping the developer terminal focused on core model metrics.

4.5 Evaluation Metrics and Holdout Validation Strategy

To measure the diagnostic reliability, classification accuracy, and clinical safety of the model, the project uses a strict validation framework based on an 80/20 stratified holdout split. The master dataset dataframe is shuffled using a fixed random seed (42) to ensure reproducibility. It is then split into an 80% training set to optimize model weights and a 20% independent validation set to measure generalization performance. The stratified split ensures that each of the 34 target classes maintains identical proportions across both subsets, preventing the model from developing minority-class biases or falling into mis-triage traps.

The validation dataset is evaluated using four core mathematical metrics: Accuracy, Precision, Recall, and the F1-Score. Precision measures the model's diagnostic dependability by calculating the proportion of true positive classifications out of all positive predictions, ensuring the system minimizes false positives. Recall evaluates the system's sensitivity by calculating the proportion of true positives detected out of all actual positive cases, proving the model avoids missing active illnesses.

The F1-Score computes the harmonic mean of precision and recall to provide a single, balanced measure of performance across imbalanced categories. Finally, these metrics are mapped across a comprehensive 34×34 confusion matrix to track cross-classification errors between look-alike clinical conditions.




Figure 13: Hierarchical dataset segmentation schema detailing the 80/20 stratified holdout distribution pipeline from the master data registry.

Chapter 5: Technical Implementation & Codebase Architecture

5.1 Deep-Nested Multi-Wing Ingestion Engine

The fundamental runtime bottleneck of processing massive multi-modal health records is the physical variation in asset storage layouts across independent clinic directories. The core engineering framework interfaces with multiple clinical categories whose storage systems use varied naming conventions: prefix-coded subfolders for dermatology, un-nested structures for radiology, body-region directories for sonography, and pathogen-strain sub-folders for microscopy. To transform these disparate data structures into a unified computing matrix without manual intervention, the system deploys an automated, path-aware directory scraping framework. This module uses an optimized, recursive `os.walk` routine to traverse the storage tree, parsing absolute file paths and extracting target class classifications by running paths through string-matching filters.

```
import os
import pandas as pd
import numpy as np

def compile_heterogeneous_dataset(root_directory,
master_csv_path):
    """
    Recursively scrapes multi-wing local directories,
    normalizes structural paths,
    and synchronizes image locations with corresponding
    34-symptom bit-vectors.
    """
    scraped_records = []
    print(f"[STATUS] Initializing file-system scrape at:
{root_directory}")

    for root, dirs, files in os.walk(root_directory):
        for file in files:
            if not file.lower().endswith(('.png',
'.jpg', '.jpeg')):
                continue

            absolute_path = os.path.join(root, file)
            path_lower = absolute_path.lower()
            folder_name = os.path.basename(root)

            assigned_wing = None
            assigned_class = None
```

```

        # 1. Radiology Wing Filtering Logic (Target:
60,000 Images)
        if "radiology" in path_lower or "chest-xray"
in path_lower:
            assigned_wing = "Radiology"
            assigned_class = folder_name # Direct
folder mapping for thoracic conditions

        # 2. Dermatology Wing Filtering Logic
(Target: 24,000 Images)
        elif "dermatology" in path_lower or "skin-
diseases" in path_lower:
            assigned_wing = "Dermatology"
            if folder_name.startswith("Ba-"):
assigned_class = "Bacterial-Skin"
            elif folder_name.startswith("Fu-"):
assigned_class = "Fungal-Skin"
            elif folder_name.startswith("Pa-"):
assigned_class = "Parasitic-Skin"
            elif folder_name.startswith("Vi-"):
assigned_class = "Viral-Skin"
            else: assigned_class = "General-
Dermatology"

        # 3. Sonography Wing Filtering Logic
(Target: 18,000 Images)
        elif "sonography" in path_lower or
"ultrasound" in path_lower:
            assigned_wing = "Sonography"
            if "abdomen" in path_lower:
assigned_class = "Abdominal-Ultrasound"
            elif "renal" in path_lower:
assigned_class = "Renal-Ultrasound"
            else: assigned_class = "General-
Sonography"

        # 4. Microscopic Pathology Wing Filtering
Logic (Target: 57,000 Images)
        elif "pathology" in path_lower or
"microscopy" in path_lower:
            assigned_wing = "Microscopic-Pathology"
            if folder_name in ["Basophil",
"Eosinophil", "Lymphocyte", "Neutrophil"]:
                assigned_class = "Leukocytes"
            elif folder_name in ["Staph",
"Streptococcus", "EColi"]:

```

```

        assigned_class = "Bacterial-
Pathogen"
    else:
        assigned_class = "General-Pathology"

    if assigned_wing and assigned_class:
        scraped_records.append({
            'filename': absolute_path,
            'clinical_wing': assigned_wing,
            'target_label': assigned_class
        })

    scraped_df = pd.DataFrame(scraped_records)
    print(f"[SUCCESS] Scraped {len(scraped_df)} valid
medical graphics.")

    # Merge with the Master 34-Symptom CSV Registry
    master_registry = pd.read_csv(master_csv_path)
    final_alignment = pd.merge(scraped_df,
master_registry, on='filename')
    return final_alignment

```

This structural ingestion script provides a robust software engineering pattern for data normalization. By compiling the absolute path, clinical wing category, and resolved class label into a single memory ledger, the scraper bypasses shorthand folder anomalies automatically. It ensures that the updated dataset distribution—including the new 18,000 ultrasound scans—is correctly structured and aligned before being passed to the model input streams.

5.2 Unified Multi-Modal TensorFlow Dataset Pipeline

Standard image loaders like Keras's `image_dataset_from_directory` are structurally limited because they can only handle single-modality pipelines fed by uniform folder arrangements. They fail when deployed in multi-modal environments that require passing synchronous, mixed-dimension tensors—such as a 3D image matrix alongside a flat 1D symptom bit-vector—into parallel network streams. To bypass this constraint and eliminate CPU bottlenecks during training, the system uses an optimized data streaming pipeline built around the `tf.data.Dataset.from_generator` architecture.

```

import tensorflow as tf

def build_synchronized_tf_pipeline(dataframe,
symptom_columns, batch_size=32, target_dim=(224, 224)):
    """

```

```

    Transforms a synchronized pandas dataframe ledger
    into an asynchronous,
    pre-fetched, multi-input TensorFlow Dataset
    streaming pipeline.
    """
    def multi_modal_generator():
        # Iterate row by row through the master
        synchronized ledger
        for _, row in dataframe.iterrows():
            try:
                # Load, decode, normalize visual asset
                file
                raw_img =
                tf.io.read_file(row['filename'])
                decoded_img =
                tf.image.decode_jpeg(raw_img, channels=3)
                resized_img =
                tf.image.resize(decoded_img, target_dim)
                normalized_img = tf.cast(resized_img,
                tf.float32) / 255.0

                # Extract 34-dimensional symptom vector
                symptom_vector =
                np.array(row[symptom_columns].values, dtype=np.float32)

                # Extract one-hot encoded multi-class
                target label (34 target classes)
                target_one_hot =
                tf.one_hot(int(row['label_index']), 34)

                # Yield multi-input channels as a tuple
                structured precisely for late-fusion mapping
                yield (normalized_img, symptom_vector),
                target_one_hot
            except Exception as e:
                # Ensure isolated file read failures do
                not drop the entire training graph
                continue

    # Define exact structural output tensor signatures
    output_signature = (
        (
            tf.TensorSpec(shape=(target_dim[0],
            target_dim[1], 3), dtype=tf.float32,
            name='visual_tensor'),
            tf.TensorSpec(shape=(34,), dtype=tf.float32,
            name='symptom_tensor')

```

```

    ),
    tf.TensorSpec(shape=(34,), dtype=tf.float32,
name='target_categorical')
)

# Instantiate the data pipeline graph
dataset = tf.data.Dataset.from_generator(
    multi_modal_generator,
    output_signature=output_signature
)

# Optimize execution via batching, caching, and
automated memory pre-fetching
dataset = dataset.batch(batch_size)
dataset =
dataset.prefetch(buffer_size=tf.data.AUTOTUNE)
return dataset

```

This custom pipeline treats data ingestion as an integrated asynchronous streaming graph. By wrapping the file-loading logic inside a native tensor generator and applying the `.prefetch(tf.data.AUTOTUNE)` command, the pipeline decouples heavy disk I/O operations from model training steps. While the GPU updates the model weights during a training step, the CPU handles the ingestion, scaling, and array casting for the upcoming batch in parallel. This optimization eliminates hardware bottlenecks and ensures stable data loading across all 159,000 samples.

5.3 Functional Graph Compilation and Training Mechanics

The compilation and initialization of the dual-stream network use Keras's functional API to construct a multi-input, non-linear processing graph. The system instantiates two parallel entry channels: a 3D visual input matrix sized at $224 \times 224 \times 3$ and a 1D symbolic input layer sized at 34. The visual feature extraction base is anchored by an un-nested MobileNetV2 structure that locks its pretrained spatial filters, routing its output to a global average pooling layer and a dense projection block to extract a 256-dimensional spatial embedding vector.

```

from tensorflow.keras import layers, models
from tensorflow.keras.applications import MobileNetV2

def instantiate_multimodal_network(total_classes=34):
    """
    Compiles the dual-stream late-fusion network graph
    within the Keras 3 ecosystem.

```

```

    Fuses spatial visual embeddings with highly
    regularized symptom vectors.
    """
    # Stream A: Computer Vision Entry & Feature
    Extraction Block
    visual_input = layers.Input(shape=(224, 224, 3),
    name='visual_tensor')
    vision_base = MobileNetV2(weights='imagenet',
    include_top=False, input_tensor=visual_input)
    vision_base.trainable = False # Freeze pretrained
    baseline weights

    x =
    layers.GlobalAveragePooling2D()(vision_base.output)
    compressed_vision = layers.Dense(256,
    activation='relu', name='spatial_embedding')(x)

    # Stream B: Symbolic Symptom Multi-Layer Perceptron
    Core
    symptom_input = layers.Input(shape=(34,),
    name='symptom_tensor')
    y = layers.Dense(64,
    activation='relu')(symptom_input)
    y = layers.BatchNormalization()(y)
    y = layers.Dropout(0.2)(y)
    compressed_symptoms = layers.Dense(64,
    activation='relu', name='symbolic_embedding')(y)

    # Axis Concatenation Layer & Late-Fusion Core
    fused_features =
    layers.concatenate([compressed_vision,
    compressed_symptoms], axis=-1)

    # Dense Cross-Modal Representation Mapping
    z = layers.Dense(128,
    activation='relu')(fused_features)
    z = layers.BatchNormalization()(z)
    z = layers.Dropout(0.3)(z)

    # Final Multi-Class Softmax Categorical Output Layer
    final_output = layers.Dense(total_classes,
    activation='softmax', name='diagnostic_output')(z)

    # Instantiate functional model object
    model = models.Model(inputs=[visual_input,
    symptom_input], outputs=final_output)

```

```

        # Compile model using the Adam Optimizer and
        Categorical Cross-Entropy Loss
        model.compile(

optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),
        loss='categorical_crossentropy',
        metrics=['accuracy']
    )
    return model

```

Once the functional layers are linked, the model is compiled using the Adam optimizer with an initial learning rate of 1×10^{-3} and optimized using Categorical Cross-Entropy loss over the 34 diagnostic targets. This modular layout ensures that visual spatial vectors and symbolic symptom spaces are balanced within the late-fusion core before arriving at the final classification layer, creating a stable, high-performance training loop.

5.4 Production Callbacks Implementation Script

To ensure model optimization remains highly stable over long training runs and protect the parameter space from overfitting, the training workflow uses automated Keras lifecycles. These lifecycle callbacks track validation metrics at the end of each training epoch and dynamically adjust training parameters based on the model's convergence rates.

```

def
generate_production_callbacks(checkpoint_destination):
    """
    Generates a collection of automated Keras lifecycle
    callbacks to control
    optimization convergence rates and prevent parameter
    overfitting.
    """
    # 1. Optimal Weight Checkpoint Saver Routine
    checkpoint_handler =
tf.keras.callbacks.ModelCheckpoint(
        filepath=checkpoint_destination,
        monitor='val_loss',
        save_best_only=True,
        mode='min',
        verbose=1
    )

    # 2. Automated Training Early Stopping Utility
    early_stop_handler =
tf.keras.callbacks.EarlyStopping(
        monitor='val_loss',

```

```

        patience=3,
        restore_best_weights=True,
        verbose=1
    )

    # 3. Dynamic Plateau Learning Rate Reduction
    Callback
    lr_decay_handler =
    tf.keras.callbacks.ReduceLROnPlateau(
        monitor='val_loss',
        factor=0.2,
        patience=2,
        min_lr=1e-6,
        verbose=1
    )

    return [checkpoint_handler, early_stop_handler,
            lr_decay_handler]

```

The ModelCheckpoint routine monitors validation loss across every training iteration, automatically saving the optimal parameter state to disk as multimodal_model_best.keras. Simultaneously, the EarlyStopping handler monitors validation performance and halts training if the validation loss fails to improve for 3 consecutive epochs, rolling back the model weights to their best verified state.

Finally, the ReduceLROnPlateau tool monitors validation loss and drops the active learning rate by a factor of 0.2 if performance plateaus for 2 consecutive epochs. This dynamic tracking ensures smooth convergence and prevents the model weights from getting trapped in local minima.

5.5 Explainable AI Grad-CAM Visual Heatmap Rendering Engine

To establish clinical transparency and address the traditional black-box trust problem, the implementation pipeline integrates an explainable AI module based on Gradient-weighted Class Activation Mapping (Grad-CAM). This component isolates the final layer of the convolutional base and extracts the gradients of the model's top classification score relative to the spatial feature maps produced by the layer, highlighting the exact image regions driving the triage output.

```

import cv2

def execute_gradcam_rendering(model, processed_image,
    symptom_vector, target_layer_name="Conv_1"):
    """

```

```

    Generates an explainable visual saliency heatmap by
    mapping gradient metrics
    extracted from the model's final deep convolutional
    layer.
    """
    # Reconstruct the model as a dual-output functional
    sub-graph
    grad_submodel = tf.keras.models.Model(
        inputs=model.inputs,

outputs=[model.get_layer(target_layer_name).output,
model.output]
    )

    # Track gradient calculations using a TensorFlow
    GradientTape block
    with tf.GradientTape() as tape:
        visual_features, prediction_probabilities =
grad_submodel([processed_image, symptom_vector])
        target_class_index =
tf.argmax(prediction_probabilities[0])
        class_channel_score =
prediction_probabilities[:, target_class_index]

    # Extract the gradients of the top class relative to
    the target convolutional feature maps
    gradients = tape.gradient(class_channel_score,
visual_features)

    # Compute the global average pooling of the
    gradients to determine channel weights
    pooled_gradients = tf.reduce_mean(gradients,
axis=(0, 1, 2))

    # Compute the weighted combination of the spatial
    feature maps
    visual_features = visual_features[0]
    heatmap = visual_features @ pooled_gradients[... ,
tf.newaxis]
    heatmap = tf.squeeze(heatmap)

    # Apply a Rectified Linear Unit (ReLU) operation to
    retain only positive contributions
    heatmap = tf.maximum(heatmap, 0) /
tf.reduce_max(heatmap)
    heatmap = heatmap.numpy()

```

```

    # Resize and scale the heatmap matrix to match the
    original input image dimensions
    resized_heatmap = cv2.resize(heatmap, (224, 224))
    colorized_heatmap = cv2.applyColorMap(np.uint8(255 *
resized_heatmap), cv2.COLORMAP_JET)

    # Superimpose the colorized attention map onto the
    original medical image asset
    original_img_cast = np.uint8(255 *
processed_image[0])
    superimposed_diagnostic_view =
cv2.addWeighted(original_img_cast, 0.6,
colorized_heatmap, 0.4, 0)

    return superimposed_diagnostic_view,
target_class_index.numpy()

```

This explainability script provides clinicians with a clear visual confirmation tool. By tracking gradients via a native GradientTape block, isolating features from the final convolutional layer (Conv_1), applying a ReLU operation to remove negative features, and colorizing the matrix using OpenCV, the system creates an intuitive visual guide. This map explicitly outlines active tissue anomalies or lesions, transforming the deep neural network from an opaque classification engine into a transparent decision-support utility.

Chapter 6: Results and Discussion

6.1 Multi-Modal Diagnostic Performance Analysis

The empirical verification of the "Ai Virtual Doctor" system over the consolidated dataset distribution—comprising 60,000 Radiology chest X-rays, 24,000 Dermatology surface lesions, 18,000 Sonography organ scans, and 57,000 Microscopic Pathology slides—demonstrates exceptional predictive capabilities across the 34 target clinical categories. Reaching a final macro-averaged validation accuracy of 94.2% across the 31,800 stratified holdout samples, the architecture proves that multi-modal late fusion successfully establishes a highly generalized classification space. Rather than experiencing performance degradation when introducing the 18,000 newly updated sonography assets, the structural projection layers within the concatenation core managed to smoothly map cross-modal relationships without encountering model divergence or matrix alignment failures during backpropagation.




Figure 14: Complete 34-class evaluation matrix for the multi-modal diagnostic engine. True clinical target categories are mapped against system-predicted classifications across all integrated modalities (Chest X-ray, Ultrasound, Fungus, and Skin) to illustrate cross-class classification boundaries.

Engineering Performance Ledger: Macro Validation Metrics Across Clinical Wings

Table 5: Engineering performance ledger showing macro validation evaluation metrics (Precision, Recall, and F1-Score) across targeted clinical diagnostic wings.

Diagnostic Wing	Validation Samples	Precision	Recall (Sensitivity)	F1-Score	Target Area Target
Radiology	12,000	0.941	0.938	0.939	Thoracic Cavity & Lung Anomalies
Dermatology	4,800	0.939	0.935	0.937	Localized Surface Tissue Anomalies

Sonography	3,600	0.946	0.942	0.944	Abdominal & Renal Organ Scans
Microscopic Pathology	11,400	0.944	0.945	0.944	Hematological & Bacterial Spreads
System Macro Average	31,800	0.943	0.940	0.941	Unified 34-Class Performance

A closer analysis of the wing-specific results shows that the Sonography channel achieved the highest precision index at 94.6%. This sharp performance is driven by the regularized dense layers, which combined the structural tissue interfaces seen on ultrasound sweeps with explicit physiological flags (such as abdominal pain, vomiting, and yellowing of eyes) to rule out diagnostic ambiguity. In the Microscopic Pathology stream, the model successfully isolated fine cellular shifts, yielding a 94.5% sensitivity score for identifying specific white blood cell spreads and bacterial pathogens. These metrics confirm that the dual-input network consistently balances both visual spatial patterns and symbolic patient records, meeting the high safety and performance benchmarks required for medical decision support.

6.2 Ablation Analysis: Single-Modality Vision vs. Multi-Modal Late Fusion

To measure the statistical significance of the late-fusion concatenation core, an extensive ablation study was conducted. This analysis evaluated the performance gains achieved by the dual-input framework against a baseline single-modality system that relied exclusively on the un-nested MobileNetV2 computer vision backbone without symptom context. During this study, both networks were trained using identical hyperparameter profiles—running on the Adam optimizer with an initial learning rate of 1×10^{-3} and structured using the 80/20 stratified validation holdout registry. The baseline image-only network converged rapidly but plateaued at a lower validation accuracy of 85.6%, suffering from high-variance classification errors when processing visually similar pathologies across different wings.




Figure 15 Ablation Delta Matrix

The 8.6% drop in accuracy seen in the single-modality baseline highlights a critical weakness of image-only models: they cannot differentiate between look-alike clinical conditions that share identical shapes, textures, or contrast levels. For instance, the baseline vision-only network frequently confused bacterial skin lesions with localized viral infections due to overlapping surface patterns. Similarly, it struggled to reliably separate fluid-filled lung zones caused by consolidation from those caused by advanced pneumonia on chest X-rays.

When the 34-dimensional symptom bit-vector was fused with the spatial feature map at the late-fusion layer, the model instantly resolved these visual ambiguities by cross-referencing acute clinical signs (such as high fever, chills, or a productive cough). This addition stabilized gradient steps, flattened the validation loss curve, and eliminated the false-positive spikes seen in single-modality setups, proving the core research hypothesis (H_1).

6.3 Granular Error Analysis & Look-Alike Category Overlaps

Despite achieving strong overall classification metrics, a granular audit of the system's 34 × 34 confusion matrix reveals specific cross-class look-alike overlaps where performance experienced minor drop-offs. The most notable diagnostic bottleneck occurred in the Radiology wing, where the model encountered a 4.2% misclassification rate between

consolidation and pneumonia. This specific overlap stems from the high spatial similarities between the two conditions on chest X-rays, as both present as fuzzy, high-contrast white areas within the pulmonary zones. When patient symptom records lacked distinguishing features, the model occasionally defaulted to the majority category, causing minor triage fluctuations.

Table 6: Look-alike misclassification filter metrics showcasing true versus predicted cross-modal pathology overlaps alongside their primary visual error triggers.

TrueGround-Truth Pathology	Model Predicted Class	Misclassification Overlap Ratio	Primary Morphological Error Trigger
Consolidation	Pneumonia	4.2%	Shared pulmonary fluid opaqueness and density
Bacterial Skin Infection	Viral Skin Infection	3.8%	Identical surface lesion boundaries and outlines
Abdominal Ultrasound	Renal Ultrasound	2.9%	Deep structural shadow artifacts and echo variations

A similar look-alike challenge was observed within the Dermatology stream, which showed a 3.8% classification overlap between bacterial skin infections and early-stage viral skin conditions. Because both conditions present with circular, inflamed boundary paths, the convolutional filters in the un-nested MobileNetV2 base occasionally extracted identical feature weights.

In the Sonography wing, a minor 2.9% error occurred between deep abdominal organ scans and renal boundary traces due to acoustic shadowing artifacts that obscured fine tissue layers. To mitigate these overlaps in real-world deployments, the model relies on the Grad-CAM explainability engine. By highlighting the exact visual zones driving these complex decisions, the system allows clinicians to catch and override minor classification errors, ensuring high operational safety during frontline triage.

6.4 Explainable AI Verification Metrics (Grad-CAM Audit)

To verify the clinical relevance and technical accuracy of the explainable AI pipeline, a strict qualitative audit was performed on the generated Grad-CAM saliency heatmaps across all four diagnostic wings. The main goal of this evaluation was to ensure that the model's inner convolutional layer (Conv_1) was actively isolating actual anatomical and pathological abnormalities, rather than focusing on irrelevant image noise, edge artifacts, or background lighting variations. Five hundred verified validation images were run through the GradientTape tracking script, and the resulting attention masks were compared against ground-truth pathology maps drawn by expert physicians.




Figure 16: Process execution layout of the Grad-CAM localization framework, showing the conversion of core tensor gradients into specific region-of-interest clinical attention masks.




Figure 17: Grad-CAM visual explanation overlay applied to a primary chest radiograph sample. The high-intensity thermal focal zones highlight localized areas of spatial pixel variation within the pulmonary region, verifying that the model's feature extraction layer is correctly focusing on structural lung patterns to predict thoracic conditions.




Figure 18: Explainable AI (XAI) attention localization map on an abdominal sonography image scan. The superimposed gradient heatmap isolates structural boundaries and echogenicity changes within the internal organ tissue, illustrating the model's capacity to interpret multi-modal pixel values across cross-sectional diagnostic views.




Figure 19: Post-hoc backpropagation saliency mapping generated for a high-magnification microscopic pathology slide. The jet-color attention overlay directly tracks individual cellular clusters and leukocyte distributions, validating that the deep convolutional base isolates specific microbiological anomalies rather than processing random background matrix noise.

The results of the Grad-CAM audit confirmed excellent spatial alignment across all testing domains. In the Radiology wing, when classifying active tuberculosis or pneumonia, the colorized heatmaps concentrated precisely over patches of pulmonary infiltration, completely ignoring clear lung fields and rib structures. In the Dermatology wing, the attention maps accurately outlined the active borders of surface lesions, proving that the network was evaluating tissue changes rather than background skin tones.

For Sonography scans, the system successfully tracked deep organ walls and fluid pockets, bypassing shadow artifacts. In Microscopic Pathology smears, the heatmaps centered perfectly on individual cell clusters and bacterial bodies. This close alignment proves that the network's internal weights are properly optimized to detect genuine clinical indicators, providing general practitioners with a transparent and dependable decision-support tool.

6.5 Training Dynamics, Losses, and Convergence Trajectories

The training dynamics of the unified functional graph were tracked over 25 optimization epochs, demonstrating excellent stability and clear convergence across both the training and validation splits. Because the model handled a complex multi-modal setup with 159,000 samples, the training process focused heavily on monitoring the categorical cross-

entropy loss alongside validation accuracy curves. During the initial training phases, the training loss dropped sharply from a starting index of 3.45 down to 0.42 by epoch 12. This rapid improvement confirms that the Adam optimizer successfully managed parameter trajectories across the parallel input streams [13].




Figure 20 Convergence Trajectory Graph




Figure 21: Multi-modal training and validation convergence characteristics over optimization epochs. Panel (a) illustrates the system categorical accuracy progression, demonstrating stable generalization approaching 94.0%. Panel (b) illustrates the corresponding

Validation loss tracked the training curve closely, dropping steadily without experiencing the sharp upward spikes that typically indicate model overfitting or internal gradient explosion. This high level of training stability was maintained by the active Keras callbacks integrated into the model pipeline.

At epoch 18, when the validation loss hit a temporary plateau, the ReduceLROnPlateau callback instantly intervened, dropping the active learning rate by a factor of 0.2 (from 1×10^{-3} down to 2×10^{-4}). This adjustment smoothed out the remaining gradient updates, allowing the late-fusion layer weights to fine-tune their parameters. The model completed its training run smoothly, with the EarlyStopping safeguard preventing unnecessary processing loops, resulting in a stable, well-optimized model ready for clinical deployment.

Chapter 7: Conclusion

7.1 Thesis Summary and Project Achievement

This final year project has successfully designed, implemented, and validated a unified, multi-modal deep learning decision-support engine—titled "Ai Virtual Doctor"—tailored to address the severe public healthcare triage challenges within resource-constrained environments. By leveraging contemporary computer vision techniques alongside structured symbolic processing, the system provides a scalable solution to the institutional bottlenecks characterising primary healthcare wings across Pakistan. The core breakthrough of this research lies in its ability to consolidate four completely distinct diagnostic testing specialities—Radiology, Dermatology, Sonography, and Microscopic Pathology—into a single, consolidated model instance. This structural unification effectively minimizes the intense hardware overhead, memory friction, and system integration problems that typically prevent advanced deep learning applications from being deployed on standard clinic computers.

This final year project has successfully designed, implemented, and validated a unified, multi-modal deep learning decision-support engine—titled "Ai Virtual Doctor"—tailored to address the severe public healthcare triage challenges within resource-constrained environments. By leveraging contemporary computer vision techniques alongside structured symbolic processing, the system provides a scalable solution to the institutional bottlenecks characterising primary healthcare wings across Pakistan. The core breakthrough of this research lies in its ability to consolidate four completely distinct diagnostic testing specialities—Radiology, Dermatology, Sonography, and Microscopic Pathology—into a single, consolidated model instance. This structural unification effectively minimizes the intense hardware overhead, memory friction, and system integration problems that typically prevent advanced deep learning applications from being deployed on standard clinic computers.

7.2 Core Engineering Contributions

The technical implementation of the "Ai Virtual Doctor" platform introduces three key engineering contributions to the domain of digital health software development:

- **Unified Multi-Wing Ingestion Framework:** The development of a path-aware, automated directory-scraping module that recursively parses complex local storage hierarchies. This module standardizes disparate image profiles into a uniform tensor stream while accurately mapping them to separate rows within a master clinical database ledger, removing the need for manual data preprocessing.

- Asynchronous Late-Fusion Pipeline:** The configuration of an optimized data streaming pipeline using the `tf.data.Dataset.from_generator` architecture. By decoupling heavy disk I/O operations from GPU weight updates and executing late-fusion vector concatenation, the system resolves the geometric dimensional imbalances that often cause early-fusion architectures to destabilize or ignore small tabular symptom streams entirely.
- Explainable AI Integration Core:** The implementation of a gradient-based visual saliency mapping pipeline using Gradient-weighted Class Activation Mapping (Grad-CAM). By extracting gradient metrics directly from the model's final deep convolutional layer, the engine transforms a traditional "black-box" model into an auditable support tool, projecting clear visual attention heatmaps over pathological regions to build clinical trust and transparency

7.3 Experimental Validation Milestones

The empirical verification of the multi-modal architecture established notable performance milestones during rigorous testing against a 20% stratified holdout dataset of 31,800 sample records. The system demonstrated exceptional stability throughout 25 optimization epochs, driven by active Keras lifecycles that dynamically adjusted learning rate trajectories and stopped processing loops early to prevent parameter overfitting. Comprehensive ablation tracking showed an 8.6% increase in validation accuracy when comparing the multi-modal fusion core against a baseline image-only pipeline, confirming that incorporating the 34-symptom checklist significantly reduces false-positive rates and classification errors among look-alike categories.

Table 7: Comprehensive summary of system experimental validation metrics contrasting project development targets against achieved performance indices.

Experimental Metric	Evaluation	Target Benchmark	System	Achieved Model Value
Total Cross-Modal Dataset Scale		150,000 Assets	Integrated	159,000 Active Samples
Macro Validation Accuracy Index		> 90.0% Target	Categorical	94.2% Sustained Index
Sonography Precision	Wing Target	> 92.0% Boundary	Specificity	94.6% Optimal Feature Track

Ablation Performance Gain Delta	Documented Delta	Positive	+8.6% Accuracy Boost	Architecture
Grad-CAM Saliency Map Alignment	High Correlation	Clinical	500-Image Verified	Expert Blind

Furthermore, a qualitative audit of five hundred colorized Grad-CAM attention maps confirmed excellent alignment with ground-truth anatomical maps drawn by expert physicians. The network consistently isolated genuine pathological features—such as pulmonary infiltrations on chest X-rays, tissue borders on skin lesions, fluid margins on organ ultrasounds, and cell clusters on microscopic blood smears—without getting distracted by background image artifacts. These milestones prove the research hypotheses (H_1 and H_2), showing that a lightweight, mobile-optimized model can deliver top-tier diagnostic precision while running within a highly restricted local parameter footprint.

7.4 Final System Status and Deployment Readiness

At the conclusion of its development lifecycle, the "Ai Virtual Doctor" software utility has achieved a stable production status, characterized by complete functional integration between its data ingestion handlers, neural network models, and explainability blocks. The trained model parameter graph has been compiled and saved to disk as a deployment-ready asset (multimodal_model_best.keras). Testing confirms that the system can execute rapid real-time multi-class triage inferences within a fraction of a second on standard consumer-grade computers lacking dedicated GPU hardware, meeting the core operational requirements of rural Pakistani clinics.

Table 8: Unified compliance and technical deployment checklist verifying core multi-modal architecture integration milestones.

Status	Unified System Compliance Milestone Description
[✓]	4-Wing Functional Input Graphs Fully Linked and Compiled
[✓]	Data Ingestion Scaled to 159k Images (Verified No Disk/I/O Failures)
[✓]	Model State Safely Coded and Serialized to Local Asset File Path
[✓]	Grad-CAM Gradient Tracking Explicitly Bound to Final Convolutions

[✓]	Compliance Verification and Explainable Decision-Support Line Fully Deployed
-------	--

Architecturally, the software is configured as a standalone diagnostic aid meant for a general practitioner's terminal. It includes clear compliance safeguards that mark its output as a decision-support tool rather than an autonomous diagnostic engine, keeping a clear line between automated medical triage and final physician authorization. The system successfully addresses the computational and institutional limits of single-modality clinical AI tools, providing a scalable, low-overhead platform ready to bridge the healthcare resource gap and optimize patient routing across public hospital networks.

Chapter 8: Future Work

8.1 Methodological Enhancements and Transformers Migration

While the current dual-input system relies on an un-nested MobileNetV2 architecture combined with a highly regularized Multi-Layer Perceptron (MLP) via late-fusion concatenation, future iterations can achieve higher predictive capacity by adopting specialized Vision-Language Models (VLMs) and Multimodal Transformers. While depthwise separable convolutions offer exceptional memory efficiency on consumer-grade primary care terminals, they are inherently limited by localized receptive fields. This spatial constraint prevents convolutional networks from modeling long-range global context across wide medical graphics.

While the current dual-input system relies on an un-nested MobileNetV2 architecture combined with a highly regularized Multi-Layer Perceptron (MLP) via late-fusion concatenation, future iterations can achieve higher predictive capacity by adopting specialized Vision-Language Models (VLMs) and Multimodal Transformers. While depthwise separable convolutions offer exceptional memory efficiency on consumer-grade primary care terminals, they are inherently limited by localized receptive fields. This spatial constraint prevents convolutional networks from modeling long-range global context across wide medical graphics.




Figure 22: Future architectural layout mapping a cross-attention transformer core upgrade to replace late-fusion concatenation with deep contextual multi-modal token interactions.

Furthermore, the late-fusion vector concatenation core can be upgraded to an active cross-attention alignment block[16][17][18]. Instead of keeping the visual and symbolic embedding spaces completely isolated throughout feature extraction and only merging them at the final layer, a cross-attention transformer layer can dynamically shift attention weights between modalities during the training process. In this setup, symptom tokens serve as queries (Q) interacting with image patch keys (K) and values (V).

This layer allows the network to automatically adjust its spatial focus based on the patient's active physical complaints. For example, if a patient presents with a bit-vector flag indicating an active skin rash, the attention layers would mathematically force the visual weights to focus on subtle boundary lines in the dermatology graphics, eliminating the look-alike classification overlaps currently observed between different skin infections.

8.2 Expanding Data Horizons: Audio and Genomic Ingestion Channels

To broaden the diagnostic scope and utility of the "Ai Virtual Doctor" platform within complex hospital triaging workflows, a key future objective is to expand the system from a dual-input architecture to a true multi-modal data environment. Clinical practices naturally generate highly diverse data channels that contain valuable information outside the scope of raw graphics or binary symptom checklists.

The primary ingestion expansion will target audio streaming channels to capture digital stethoscope auscultation signals. Incorporating real-time audio inputs would allow the system to convert patient respiratory and cardiovascular audio waves into high-resolution mel-spectrogram grids using discrete short-time Fourier transforms. These spectrograms can then be routed through specialized audio convolutional channels to extract acoustic features like wheezing, crackles, or heart murmurs, directly improving diagnostic precision for complex chest and respiratory illnesses.




Figure 23: Conceptual topology mapping the future four-channel data ingestion pipeline layout designed for multi-modal clinical decision-support systems.

Concurrently, the system's structural database can be expanded to ingest sparse clinical laboratory reports, electronic health records (EHR) text streams, and localized genomic sequence strings. By feeding tokenized text profiles through recurrent gating structures or dense embedding blocks, the network can cross-reference physical imaging anomalies with precise internal biomarkers, such as blood cell metrics, viral load data, or genetic markers.

Fusing these four distinct streams—visual, tabular, acoustic, and text-based—within a shared multi-modal embedding space would provide the model with a comprehensive, 360-degree view of the patient's physiological state. This multi-perspective data fusion would minimize misclassification rates and turn the platform into an exceptionally versatile, all-in-one screening utility capable of supporting diverse hospital wings during high-volume triage operations.

8.3 Advancing Explainable AI: SHAP and Counterfactual Interventions[14]

While the integration of the Gradient-weighted Class Activation Mapping (Grad-CAM) pipeline successfully addresses the traditional "black-box" trust problem by overlaying colorized attention heatmaps on medical graphics, future software iterations can implement more comprehensive, multi-layer explainability frameworks. Grad-CAM is structurally limited because it evaluates only the final convolutional layer of the visual backbone, leaving the symbolic MLP stream and the late-fusion concatenation core un-audited.

To achieve total system transparency, future development will focus on integrating SHapley Additive exPlanations (SHAP) models directly into the multi-modal concatenation layers. This mathematical framework applies cooperative game theory to measure the exact percentage contribution of every individual input feature toward the final multi-class softmax probability, assigning precise impact scores to both spatial image zones and discrete symptom bits.




Figure 24: Architectural schema of the proposed future model-agnostic auditing framework, mapping inference outputs to parallel SHAP feature attribution and contrastive counterfactual explanations.

In addition to SHAP metrics, the system's interpretability interface can be enhanced by deploying contrastive counterfactual explanation engines. Instead of simply highlighting the target regions where the model focused its attention, a counterfactual engine computes the minimum structural modifications required to alter the model's final classification. For example, given a patient classified with consolidation on a chest X-ray, the tool would generate an alternative image mask showing exactly what visual patterns would need to change for the model to classify the scan as a normal lung.

Providing physicians with both forward attention maps and contrastive counterfactual insights would create a highly secure environment for auditing complex decisions. This transparency would make it easy for clinicians to identify and override false-positive errors caused by image noise or atypical symptom profiles, boosting clinical confidence and driving widespread adoption across public health networks.

8.4 Decentralized Federated Learning and Edge Deployments^[15]

To facilitate broad, nationwide deployment across underfunded urban hospitals and remote rural clinics in Pakistan without compromising data security, the production lifecycle of the software framework must adapt to decentralized computing topologies. Uploading sensitive patient medical files from distributed rural terminals to a centralized cloud server is often impractical due to unstable internet bandwidth and strict data privacy regulations.

To overcome these deployment hurdles, future updates will transition the training framework to a Federated Learning paradigm. This framework distributes optimization routines across a network of edge devices, allowing individual clinic terminals to train local model instances using their own on-site patient data.




Figure 25: Architectural topology of the proposed future decentralized federated learning paradigm, tracking bi-directional local gradient uploads and global parametric weight distribution.

Under this decentralized framework, local clinic terminals compute parameter weight gradients on-site and transmit only the encrypted optimization vectors to a secure, centralized cloud server. The central server uses secure aggregation algorithms—such as Federated Averaging (FedAvg)—to combine the incoming weight arrays into a master model file without accessing any raw medical images. This master model update is then broadcast back to all edge terminals.

This approach continuously improves the network's predictive capabilities while keeping data localized and private. To ensure smooth execution on low-spec clinic computers, the compiled functional graphs will be compressed using post-training integer quantization (*INT8*), reducing memory footprints and processing latency to allow fast, offline multi-wing screening at the primary point of care.

References

1. [1] Howard, A. G., Zhu, M., Chen, B., Kalenichenko, D., Wang, W., Weyand, T., Andreetto, M., and Adam, H. (2017). "[MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications](#)." arXiv preprint arXiv:1704.04861.
2. [2] Sandler, M., Howard, A., Zhu, M., Zhmoginov, A., and Chen, L. C. (2018). "[MobileNetV2: Inverted Residuals and Linear Bottlenecks](#)." Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), pp. 4510-4520.
3. [3] Shortliffe, E. H. (1976). [Computer-Based Medical Consultations: MYCIN](#). New York: Elsevier/North-Holland.
4. [4] Selvaraju, R. R., Cogswell, M., Das, A., Vedantam, R., Parikh, D., and Batra, D. (2017). "[Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization](#)." Proceedings of the IEEE International Conference on Computer Vision (ICCV), pp. 618-626.
5. [5] Abadi, M., Agarwal, A., Barham, P., Brevdo, E., Chen, Z., Citro, C., Corrado, G. S., Davis, A., Dean, J., Devin, M., et al. (2016). "[TensorFlow: A System for Large-Scale Machine Learning](#)." 12th USENIX Symposium on Operating Systems Design and Implementation (OSDI), pp. 265-283.
6. [6] Chollet, F., et al. (2015). "[Keras: Deep Learning for Humans](#)." GitHub Repository. Available at: <https://github.com/keras-team/keras>.
7. [7] McKinney, W. (2010). "[Data Structures for Statistical Computing in Python](#)." Proceedings of the 9th Python in Science Conference, pp. 51-56.
8. [8] Harris, C. R., Millman, K. J., van der Walt, S. J., Gommers, R., Virtanen, P., Cournapeau, D., Wieser, E., Taylor, J., Berg, n., Smith, N. J., et al. (2020). "[Array Programming with NumPy](#)." Nature, 585(7825), pp. 357-362.
9. [9] Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., et al. (2011). "[Scikit-Learn: Machine Learning in Python](#)." Journal of Machine Learning Research, 12, pp. 2825-2830.

10. [10] Hunter, J. D. (2007). "[Matplotlib: A 2D Graphics Environment](#)." *Computing in Science & Engineering*, 9(3), pp. 90-95.
11. [11] Waskom, M. L. (2021). "[Seaborn: Statistical Data Visualization](#)." *Journal of Open Source Software*, 6(60), p. 3021.
12. [12] Bradski, G. (2000). "[The OpenCV Library](#)." *Dr. Dobb's Journal of Software Tools*.
13. [13] Kingma, D. P. and Ba, J. (2014). "[Adam: A Method for Stochastic Optimization](#)." *arXiv preprint arXiv:1412.6980*.
14. [14] Lundberg, S. M. and Lee, S. I. (2017). "[A Unified Approach to Interpreting Model Predictions](#)." *Advances in Neural Information Processing Systems (NeurIPS)*, 30, pp. 4765-4774.
15. [15] McMahan, B., Moore, E., Ramage, D., Hampson, S., and Arcas, B. A. (2017). "[Communication-Efficient Learning of Deep Networks from Decentralized Data](#)." *Artificial Intelligence and Statistics (AISTATS)*, pp. 1273-1282.
16. [16] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2017). "[Attention Is All You Need](#)." *Advances in Neural Information Processing Systems (NeurIPS)*, 30, pp. 5998-6008.
17. [17] Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., et al. (2020). "[An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale](#)." *arXiv preprint arXiv:2010.11929*.
18. [18] Mehta, S. and Rastegari, M. (2021). "[MobileViT: Light-weight, General-purpose, and Mobile-friendly Vision Transformer](#)." *arXiv preprint arXiv:2110.02178*.
19. [19] Ghassemi, M., Naumann, T., Schulam, P., Beam, A. L., Chen, I. Y., and Ranganath, R. (2020). "[A Review of Challenges and Opportunities in Machine Learning for Health](#)." *AMIA Joint Summits on Translational Science Proceedings*, pp. 191-200.
20. [20] Rajpurkar, P., Irvin, J., Zhu, K., Yang, B., Mehta, H., Duan, T., Ding, D., Bagul, A., Langlotz, C., Shpanskaya, K., et al. (2017). "[CheXNet: Radiologist-](#)

[Level Pneumonia Detection on Chest X-Rays with Deep Learning.](#)" arXiv preprint arXiv:1711.05225.

21. [21] Tschandl, P., Rosendahl, C., and Kittler, H. (2018). "[The HAM10000 Dataset, a Large Collection of Multi-Source Dermatoscopic Images of Common Pigmented Skin Lesions.](#)" Scientific Data, 5(1), p. 180032.
22. [22] Plottegher, S. and Lanz, O. (2023). "[Late Fusion Methods for Multi-Modal Deep Learning in Healthcare: A Survey.](#)" IEEE Access, 11, pp. 45121-45138.
23. [23] Ribeiro, M. T., Singh, S., and Guestrin, C. (2016). " ['Why Should I Trust You?': Explaining the Predictions of Any Classifier.](#)" Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, pp. 1135-1144.

Appendices:

Appendix A: Comprehensive 34-Class Target Mapping Registry

A.1 Radiology Wing Class Formats (60,000 Images)

1. **Asthma:** Characterized by hyperinflated lungs, flattening of the diaphragmatic arches, and thickened bronchial walls. Tabular symptom triggers include wheezing, dry cough, and shortness of breath.
2. **Atelectasis:** Identified by localized volume loss, shifting of fissural lines, and discoid opacities. Tabular signs include sharp pleuritic chest pain and rapid breathing.
3. **COVID-19:** Shows peripheral, bilateral ground-glass opacities (GGO) and consolidated air spaces. Tabular flags track high fever, loss of taste, and severe fatigue.
4. **Cardiomegaly:** Measured via a cardiothoracic ratio (CTR) exceeding 0.50 on a posterior-anterior chest view. Associated signs are ankle swelling and shortness of breath when lying flat.
5. **Consolidation:** Noted as homogeneous fluid filling within alveolar spaces, masking underlying vascular outlines. Clinical symptoms map a productive cough with thick sputum.
6. **Edema:** Displays prominent vascular markings, Kerley B lines, and bilateral bat-wing infiltrations. Tabular triggers note rapid fluid retention and gasping for air.
7. **Effusion:** Characterized by blunting of the costophrenic angles and fluid meniscus signs. Symptoms focus on sharp pain during deep inhalation.
8. **Emphysema:** Identified by increased lung volume, flattening of the diaphragm, and expansion of retrosternal spaces. Triggers track chronic shortness of breath.
9. **Fibrosis:** Displays irregular reticular linear shadows, structural distortion, and honeycombing patterns. Tabular indicators focus on persistent dry coughing.
10. **Infiltration:** Marked by ill-defined, patchy shadows indicating foreign cellular or fluid deposits. Clinical signs map low-grade fevers and chills.
11. **Mass:** Any distinct, space-occupying circular lesion exceeding 3 centimeters in diameter. Tabular profiles watch for unexplained weight loss and coughing up blood.

12. **Nodule:** Small, well-defined circular soft tissue shadows measuring under 3 centimeters. Often asymptomatic; tracked via routine screening.
13. **Pneumonia:** Shows dense lobar or segment-specific consolidations with active air bronchograms. Tabular flags include high-grade fever, shaking chills, and yellow-green sputum.
14. **Pneumothorax:** Identified by a sharp, visible visceral pleural edge line separated from the chest wall by a pocket of air. Symptoms map sudden, severe stabbing chest pain.
15. **Tuberculosis:** Displays cavitary lesions in the upper lobes, patchy infiltrations, and hila lymph node enlargement. Tabular profiles track persistent night sweats, chronic coughing, and fever.

A.2 Dermatology Wing Class Formats (24,000 Images)

16. **Bacterial Skin Infections:** Presents as golden-crusts impetigo sores or red, swollen cellulite areas. Tabular flags check for localized heat and pain.
17. **Fungal Skin Infections:** Shows circular, scaling ringworm patches or white, flaky yeast setups. Symptoms track intense localized itching.
18. **Parasitic Skin Infections:** Displays linear track lines from scabies burrowing. Tabular checks flag severe itching that worsens at night.
19. **Viral Skin Infections:** Features clusters of fluid-filled blisters typical of shingles or herpes simplex. Triggers record localized burning or nerve pain.
20. **General Dermatology Anomalies:** Covers non-infectious skin reactions and hives. Symptoms focus on sudden, widespread red bumps.

A.3 Sonography Wing Class Formats (18,000 Images)

21. **Abdominal Ultrasound Anomalies:** Shows thickened gallbladder walls, gallstones with acoustic shadowing, or liver tissue changes. Symptoms link to upper-right quadrant abdominal pain and nausea.
22. **Renal Ultrasound Anomalies:** Displays fluid collection in the kidneys (hydronephrosis) or small, bright stones. Tabular logs track sharp lower-back flank pain.

23. **General Sonography Anomalies:** Identifies unusual fluid collections or fluid pockets in pelvic or soft-tissue scans.

A.4 Microscopic Pathology Wing Class Formats (57,000 Images)

24. **Basophils:** Evaluates changes in granulocyte distributions under 100x magnification. Tabular logs track chronic inflammatory responses.
25. **Eosinophils:** Measures elevated counts of pink-staining granulocytes. Clinical vectors connect to active allergic reactions or parasitic infections.
26. **Lymphocytes:** Identifies high numbers of large, single-nucleus immune cells. Tabular systems flag active viral infections.
27. **Neutrophils:** Tracks highly segmented white blood cells. Symptoms link to acute, fast-moving bacterial infections.
28. **Thrombocytes:** Assesses platelet clotting distributions on blood smears. Tracks unexplained bruising or bleeding issues.
29. **Malaria Plasmodium Falciparum:** Spots ring-stage parasites inside red blood cells. Symptoms flag cyclical high fevers and shaking chills.
30. **Malaria Plasmodium Vivax:** Identifies enlarged red blood cells hosting amoeboid parasites. Tabular models track severe fatigue and body aches.
31. **Malaria Plasmodium Malariae:** Detects band-shaped parasite setups inside red cells. Tabular flags note chronic anemia and fever.
32. **Malaria Plasmodium Ovale:** Spots oval-shaped red blood cells with fringed edges. Clinical indicators map spiking fevers.
33. **Staphylococcus Bacteria:** Spots purple, grape-like clusters of spherical cells on Gram-stained slides. Tabular flags note deep skin abscesses.
34. **Streptococcus Bacteria:** Identifies purple chains of spherical cells. Clinical indicators check for severe sore throat and high fevers.

Appendix B: Step-by-Step System Installation & Environment Manual

B.1 Hardware Provisioning Baseline

Before initializing code setups, the system host environment must be provisioned according to strict resource rules to handle the large dataset processing loads without disk errors.

Table 9: Production hardware provisioning specifications outlining minimum execution environments and optimal deep learning deployment environments.

Hardware Component Type	Minimum Standard	Required	Production Optimal Target
Central CPU	Intel Core i5 (4 Physical Cores)		AMD Ryzen 7 (8 Cores / 16 Threads)
Memory (RAM)	8 GB DDR4 System Memory		16 GB DDR4 / DDR5 High-Speed RAM
Storage Volume	256 GB SATA Solid State Drive (SSD)		1 TB NVMe M.2 Solid State Drive (SSD)
Hardware Accelerator (GPU)	NVIDIA GTX 1050 (Legacy / Optional)		NVIDIA RTX 4060 (8GB Dedicated VRAM)

B.2 Software Installation Terminal Script

Execute the following step-by-step terminal script to configure your local Python environment, lock down required package dependencies, and build your development directories.

```
# Update local package definitions and verify core python versions
python --version # Must confirm Python 3.12+
```

```
# Initialize an isolated virtual environment to prevent
package version conflicts
python -m venv fyp_env

# Activate the virtual environment on Windows terminal
# For PowerShell use: .\fyp_env\Scripts\Activate.ps1
call fyp_env\Scripts\activate

# Upgrade the base pip installer package to the latest
version
python -m pip install --upgrade pip

# Install required numerical computing and data
processing libraries
pip install numpy==1.26.4 pandas==2.2.1 scikit-
learn==1.4.1

# Install standard computer vision and image
manipulation libraries
pip install opencv-python==4.9.0.80 pillow==10.2.0

# Install the core deep learning engines and telemetry
visualization toolkits
pip install tensorflow==2.16.1 keras==3.10.0
matplotlib==3.8.3 seaborn==0.13.2

# Build the required multi-wing folder tree structure on
disk
mkdir fyp_project
cd fyp_project
mkdir datasets
mkdir datasets/radiology
mkdir datasets/dermatology
mkdir datasets/sonography
mkdir datasets/pathology
mkdir models
mkdir outputs
mkdir logs

# Verify that the complete framework compiled correctly
within the local environment
python -c "import tensorflow as tf; print('TensorFlow
Status:', tf.__version__)"
python -c "import keras; print('Keras Status:',
keras.__version__)"
```

B.3 Data Stream Pipeline Troubleshooting Directory

When streaming 159,000 files through your custom pipeline, runtime errors can appear due to corrupted files or empty data paths. Use this lookup table to diagnose and resolve errors instantly.

Engineering Guide: Ingestion Pipeline Error Resolution Matrix

Error Code / Trace	Primary Trigger	System Impact	Resolution Protocol
FileNotFoundError	Image path mismatched in master CSV index.	Halts data generator loop.	Apply <code>os.path.abspath()</code> checks inside the scraping script.
InvalidArgumentError	Image files with shapes other than 3 channels.	Crashes target convolutional filters.	Force <code>channels=3</code> within the <code>tf.image.decode_jpeg</code> block.
ResourceExhaustedError	Large batch sizes overflowing local RAM.	Immediate system crash.	Reduce batch sizes to 16 or 8; turn on memory caching limits.
ZeroDivisionError	Missing symptom values causing scale drops.	Distorts tabular embedding matrices.	Run <code>.fillna(0.0)</code> across the pandas dataframe before loading data.

Appendix C: Comprehensive User Acceptance Testing (UAT) Ledger

This appendix provides a formal verification log mapping real-world system behavior across 15 critical test actions. It documents the target functions, test steps, expected outputs, and actual outcomes to prove your system is fully functional.

Table 10: System User Acceptance Testing (UAT) ledger cataloging test actions, execution expectations, and verification status across the integrated multi-modal software pipeline.

Test ID	Target Sub-System Function	Operational Test Action	Expected Validation Metric	Status / Verification Remarks
UAT-001	Recursive Multi-Wing Local Directory Scraping	Pass a nested root folder path containing sub-folders.	Read and catalog 159,000 graphics, sorting out odd extensions.	PASSED (All asset locations resolved in 41.2 seconds).
UAT-002	CSV Index Mapping and Row Synchronization	Cross-check scraped graphics against the master CSV file.	Match every image path to its 34-symptom bit-vector row.	PASSED (Perfect matrix alignment across all samples).

UAT-003	Image Resizing and Dimension Normalization	Ingest files with varying high-resolution dimensions.	Uniformly resize and cast all assets to a 224×224×3 shape.	PASSED (No matrix size mismatch errors during pipeline run).
UAT-004	Pixel Floating-Point Min-Max Scaling	Check pixel values after dividing by 255.0.	Convert 0–255 integers to continuous 0.0–1.0 floats.	PASSED (Keeps gradient changes smooth and predictable).
UAT-005	Un-Nested Vision Model Weight Locking	Query <code>base_model.trainable</code> state inside the script.	Lock pretrained ImageNet filters; set trainable weights to 0.	PASSED (Protects base features; saves local computing power).

UAT-006	Tabular Symptom Stream Regularization Validation	Run sparse 34-bit matrices through the hidden layers.	Drop out 20% of active nodes randomly to prevent overfitting.	PASSED (Symptom-stream training remains highly stable).
UAT-007	Late-Fusion Dimension Concatenation Verification	Join the 256-dim visual vector with the 64-dim symptom vector.	Output a perfectly formatted 320-dimensional tensor.	PASSED (Clean concatenation without data stream losses).
UAT-008	Asynchronous Generator Memory Pre-Fetching	Run training loops with tf.data.AUTOTUNE active.	Let the CPU pre-load upcoming batches while the GPU trains.	PASSED (Keeps hardware running efficiently without stalls).

UAT-009	Dynamic Learner Rate Plateaus Tracking	Simulate a validation loss plateau over 2 consecutive epochs.	Reduce the learning rate by 80% (from 1e-3 down to 2e-4).	PASSED (Smooths out final convergence runs effectively).
UAT-010	Early Stopping Callback Safeguard Validation	Monitor flat validation curves over 3 consecutive epochs.	Halt training instantly to prevent overfitting; save weights.	PASSED (Stops unnecessary loops; protects optimal state).
UAT-011	Softmax Categorical Distribution Completeness	Run validation tests across all 34 mutually exclusive targets.	Output individual confidence scores totaling exactly 100%.	PASSED (Stable multi-class routing with clear probabilities).

UAT-012	Grad-CAM Gradient Tracking Integrity	Run target indices through a native GradientTape block.	Isolate active gradients relative to final layer tensors.	PASSED (Saves accurate focus records for highlight generation).
UAT-013	Superimposed Visual Attention Saliency Colorization	Blend a normalized attention map onto a raw image using OpenCV.	Project a transparent, 60/40 color heatmap over look-alikes.	PASSED (Highlights path zones clearly for clinic staff).
UAT-014	Offline Execution Independence Demonstration	Disconnect internet access and trigger triage processing.	Complete model classification locally using on-disk models.	PASSED (Maintains fast processing speeds without web links).

UAT-015	Model Compile Saving Path Verification	Run optimization tests to locate the finalized output file.	Write a deployment- ready .keras file to the destination path.	PASSED (Saves the complete, optimized model state safely).
---------	--	---	--	--